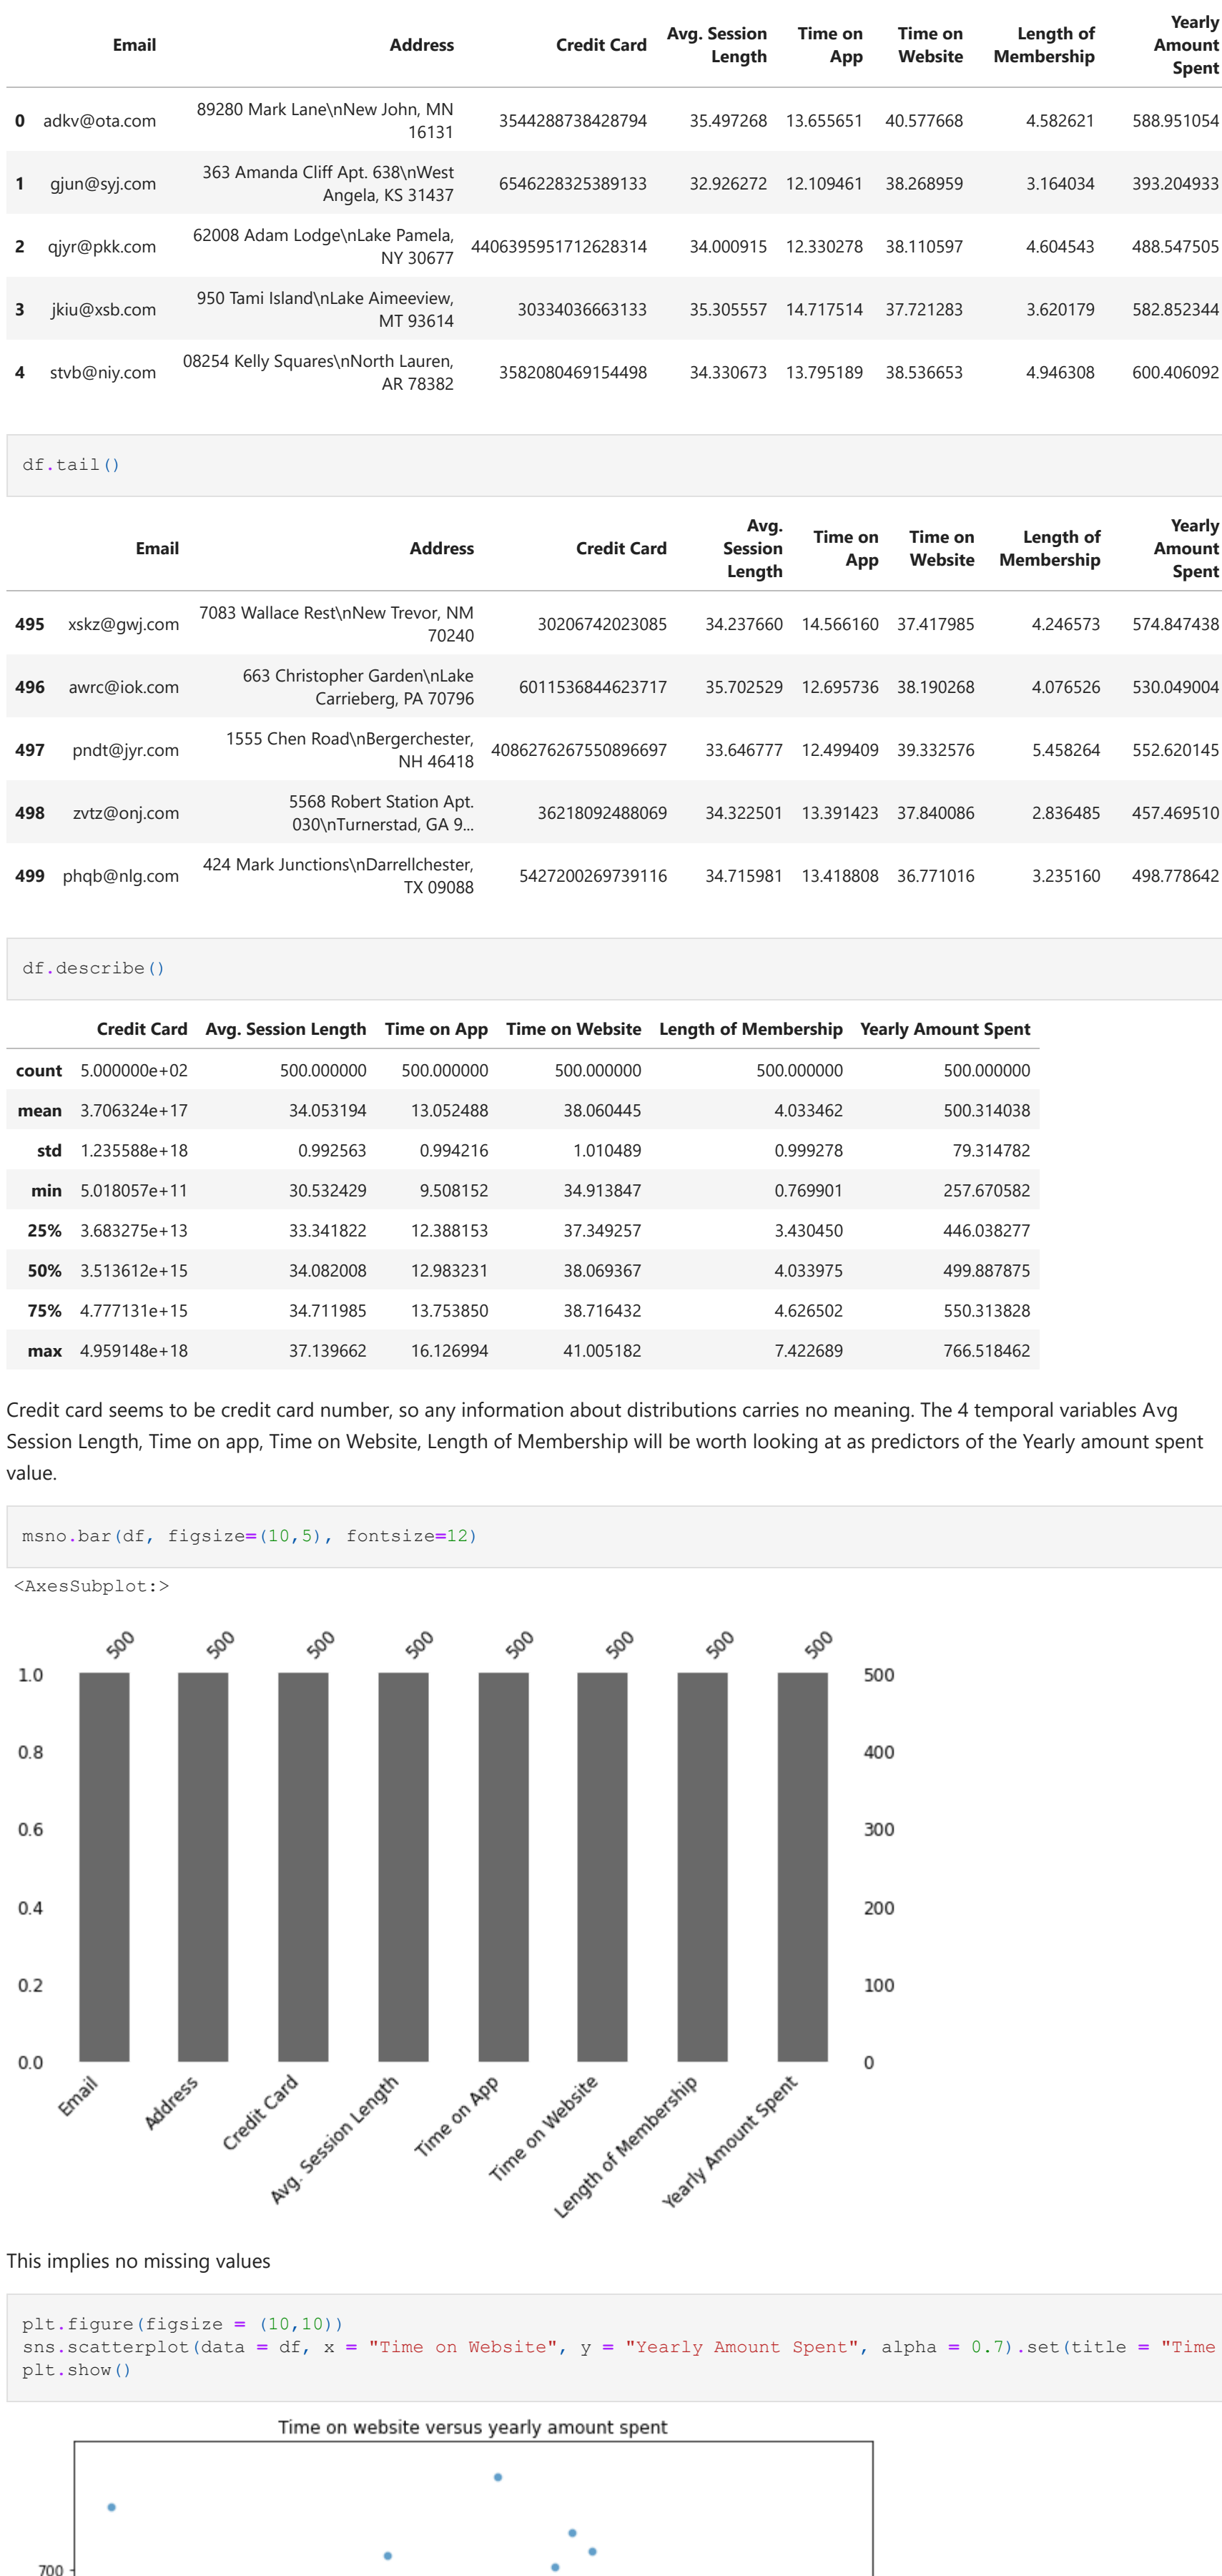


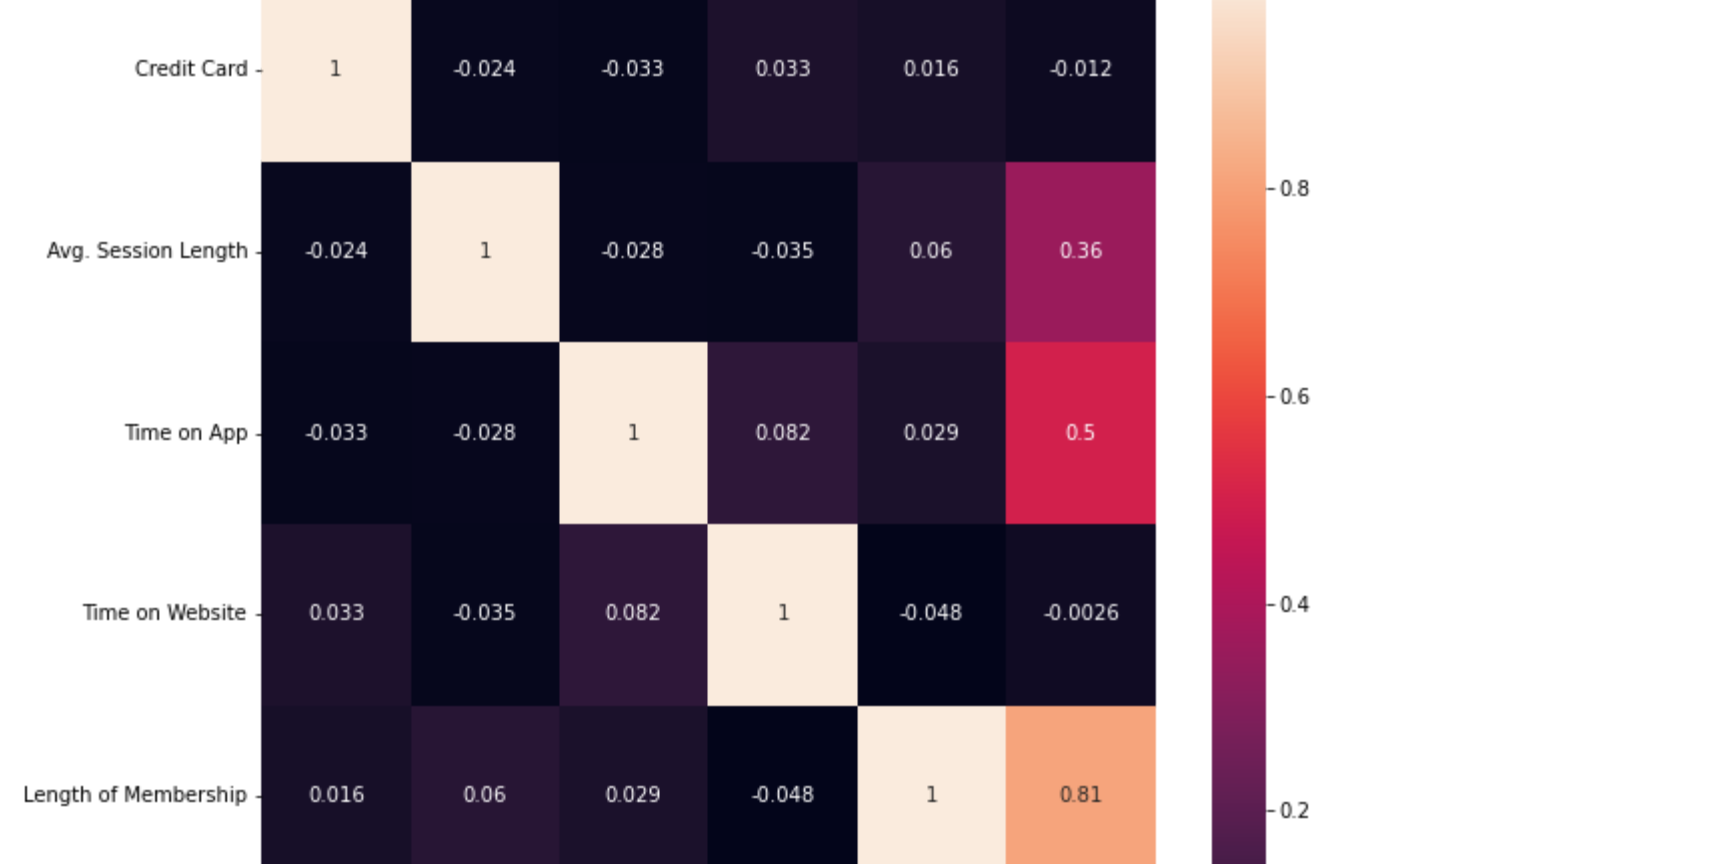
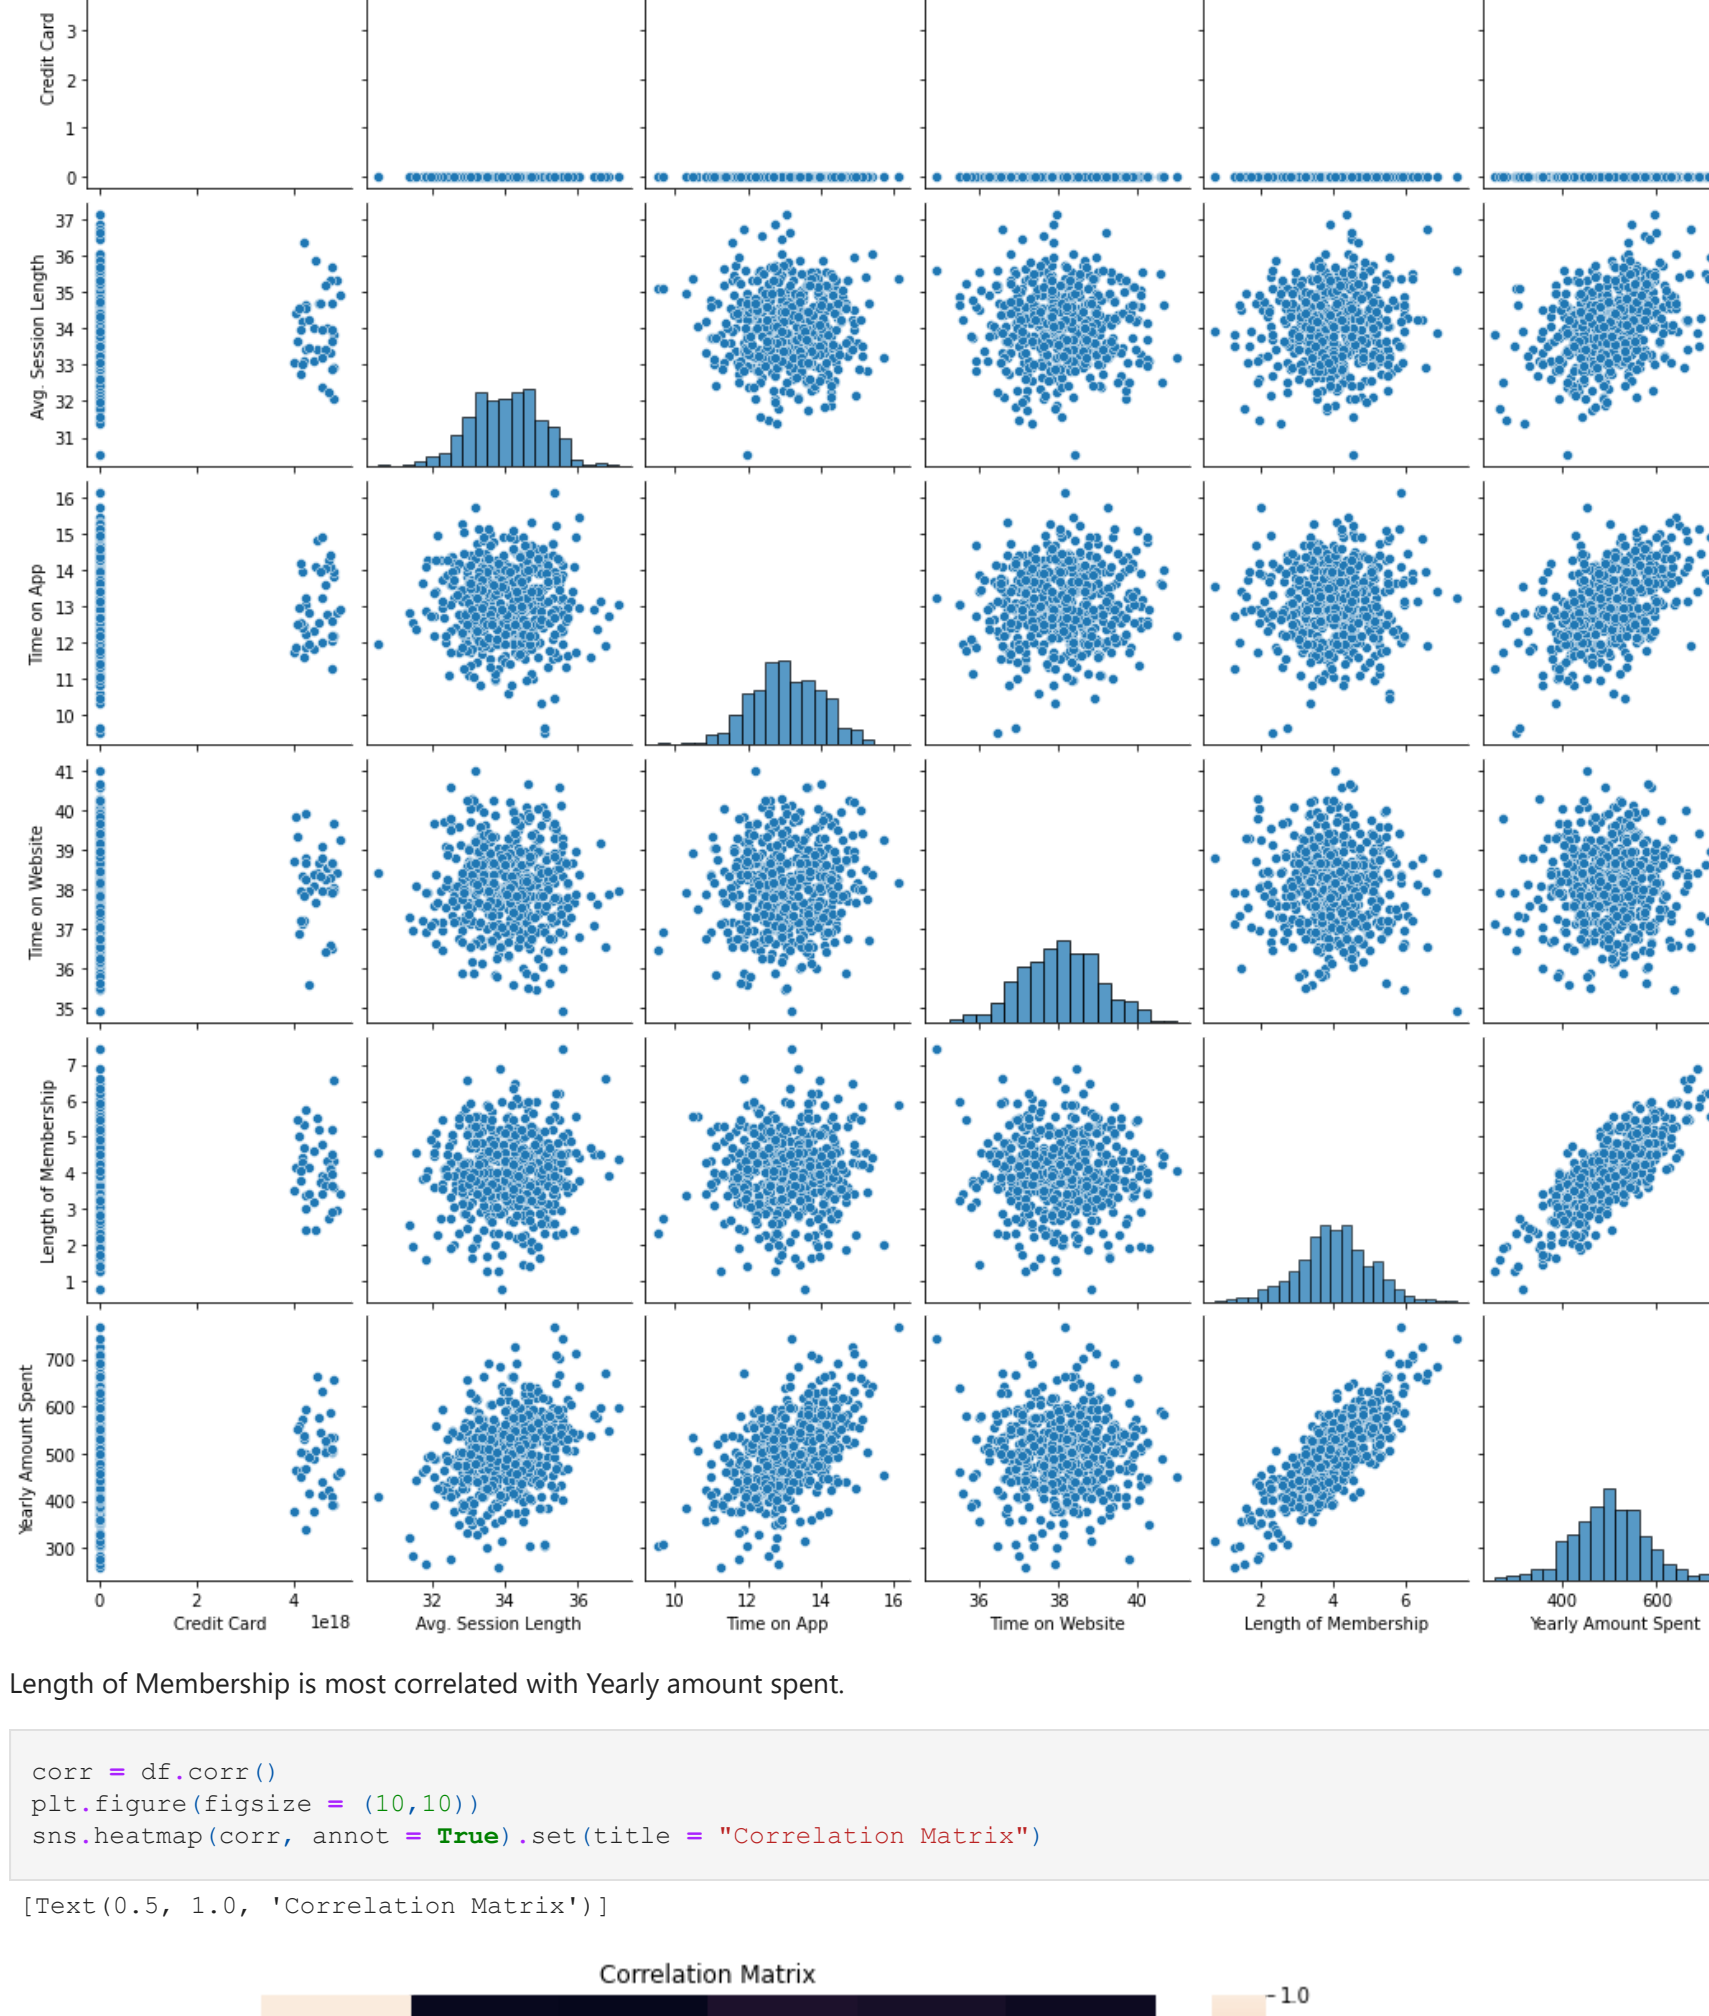
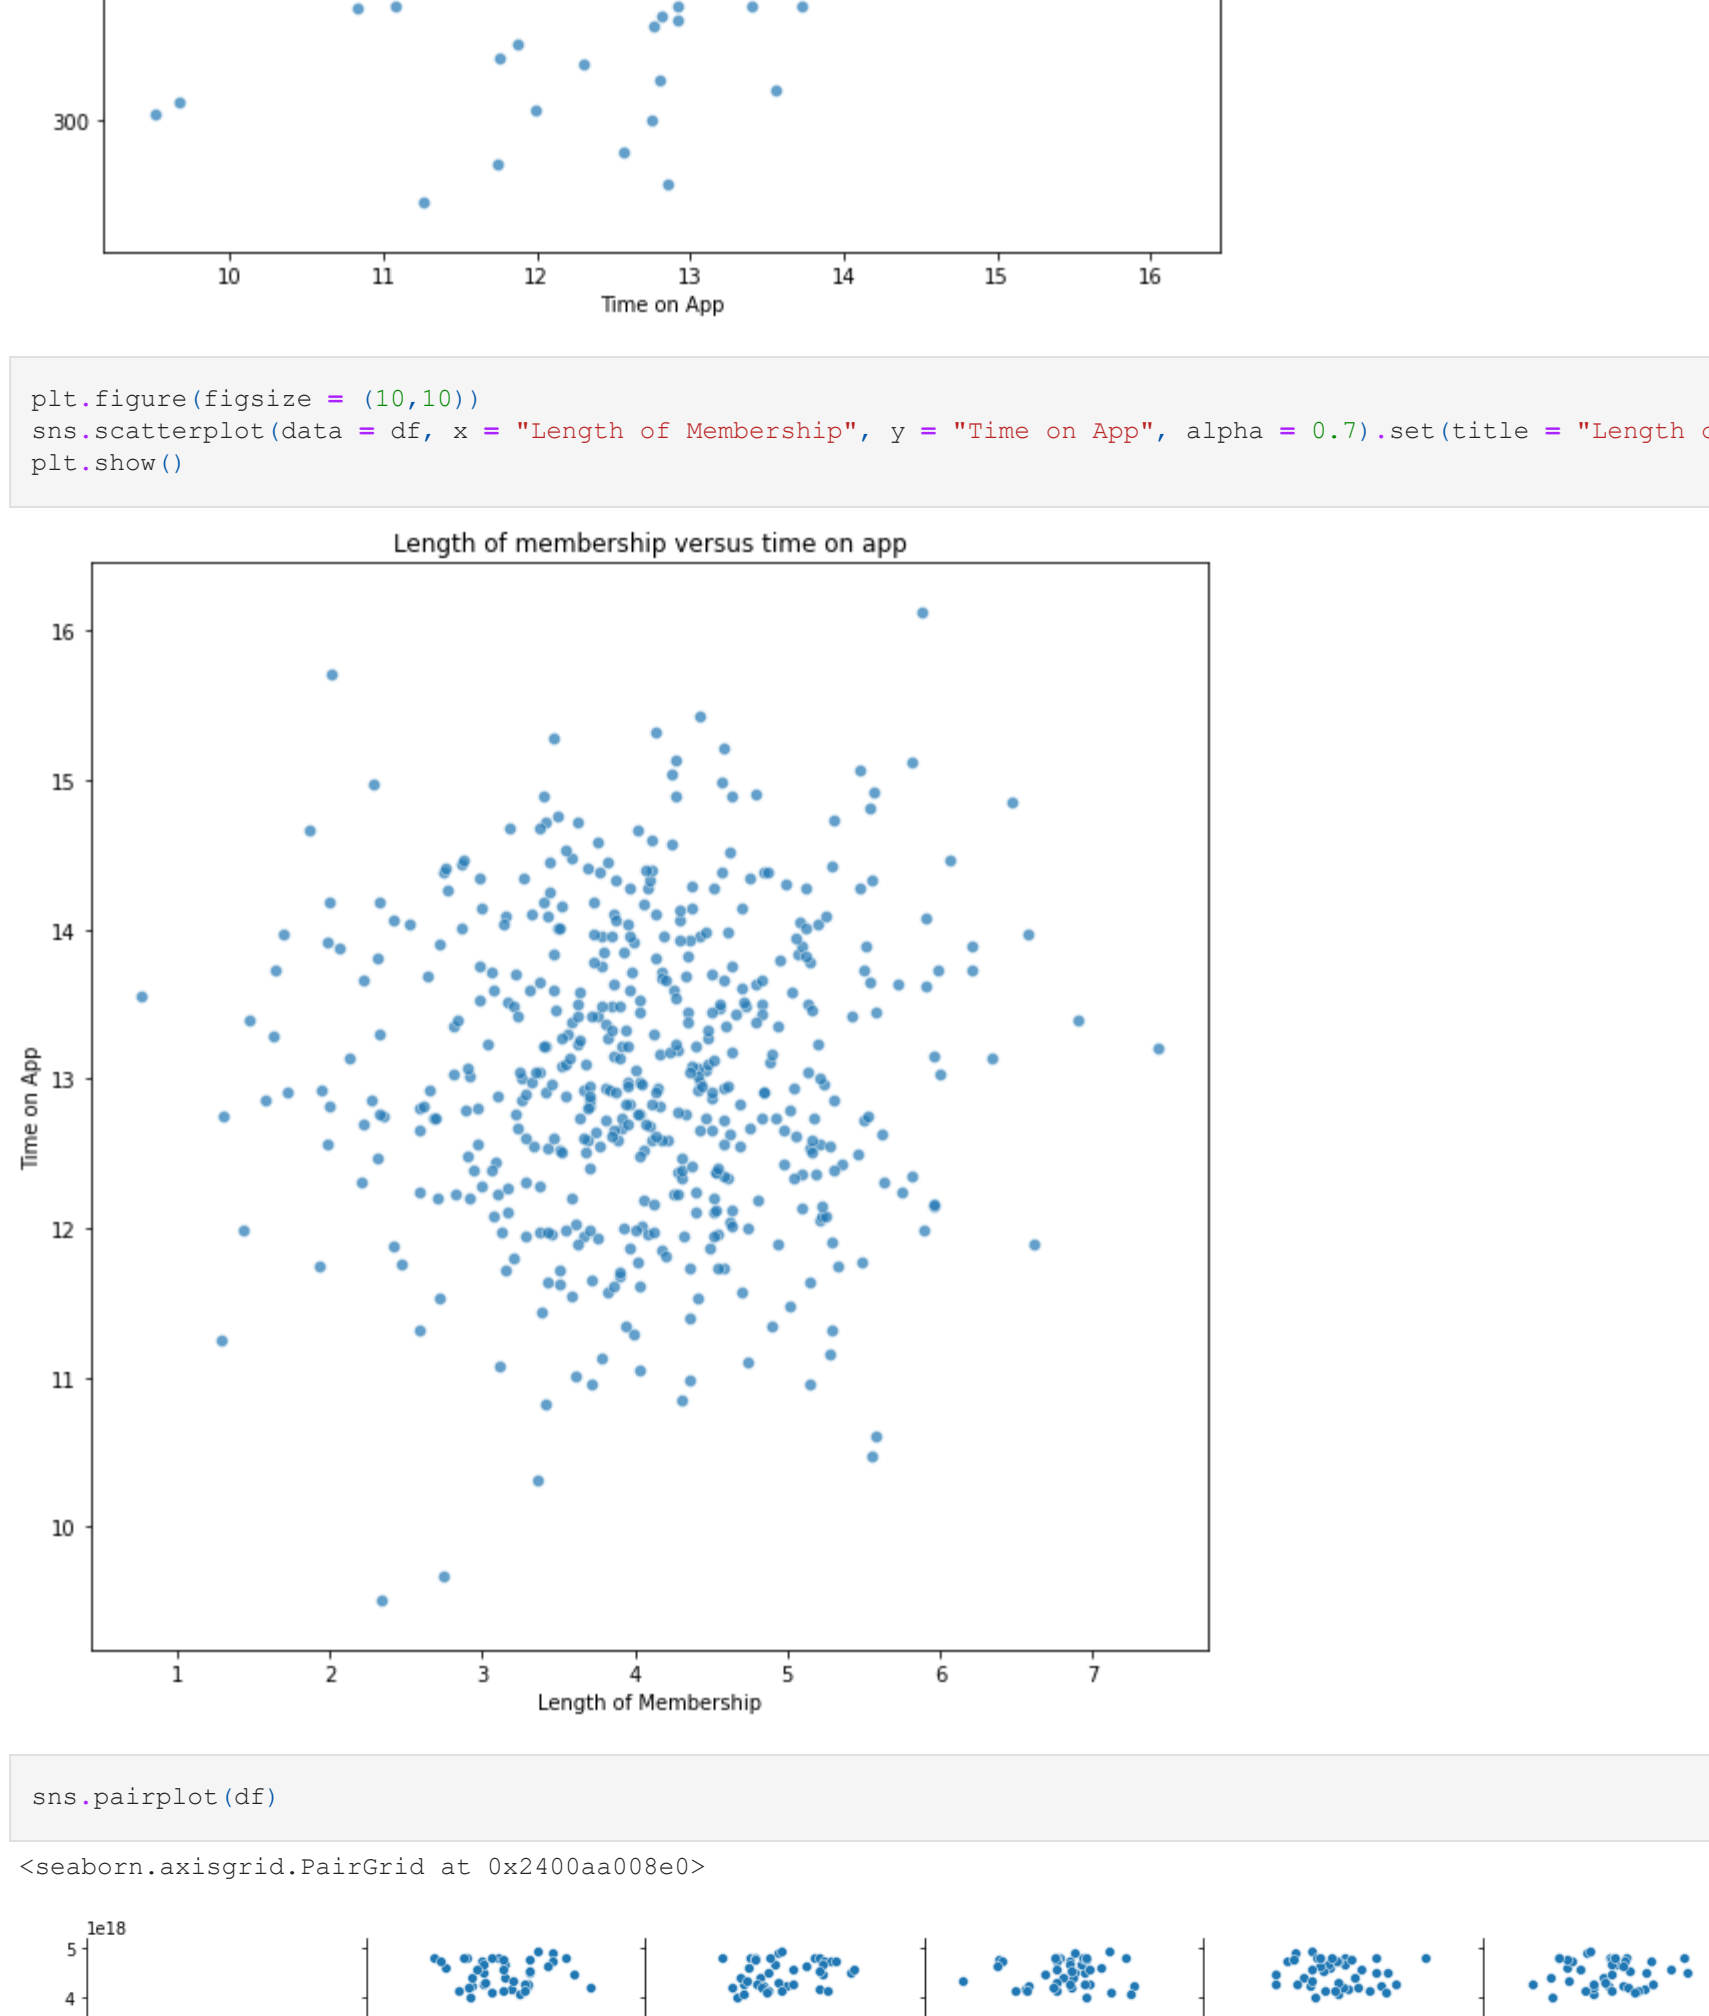
This jupyter notebook is prepared by Matthias Rathbun.

1. Load Data and perform general EDA

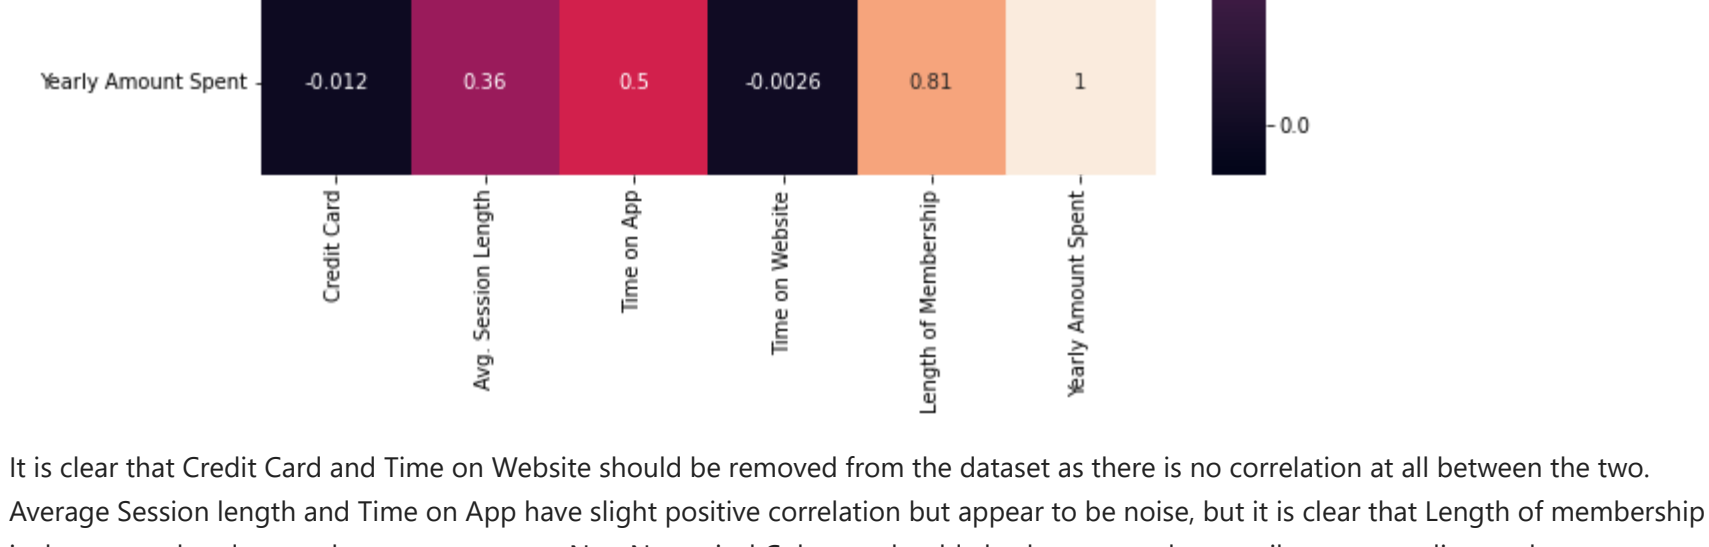
1. Import libraries: pandas, numpy, matplotlib (set %matplotlib inline), matplotlib.pyplot, seaborn, missingno, scipy's stats, sklearn (1 pt)
2. Import the data to a dataframe and show the count of rows and columns (1 pt)
3. Show the top 5 and last 5 rows (1 pt)
 - A. Explain in words about the description of any two variables (1 pt)
5. Show any missing value analysis (1 pt)
6. Plot various scatter plots to understand the data:
 - A. Yearly amount Spent vs Time on Website
 - B. Yearly amount Spent vs Time on App
 - C. Length of membership vs Time on App
- E. Also, plot sns heatmap based on correlation with Anno=True and discuss which columns must be removed based on that and which column is mostly interesting and related to Yearly Amount Spent?
- F. Generate a scatter plot with the interesting column you found in the last step against the Yearly Amount Spent



This implies no missing values

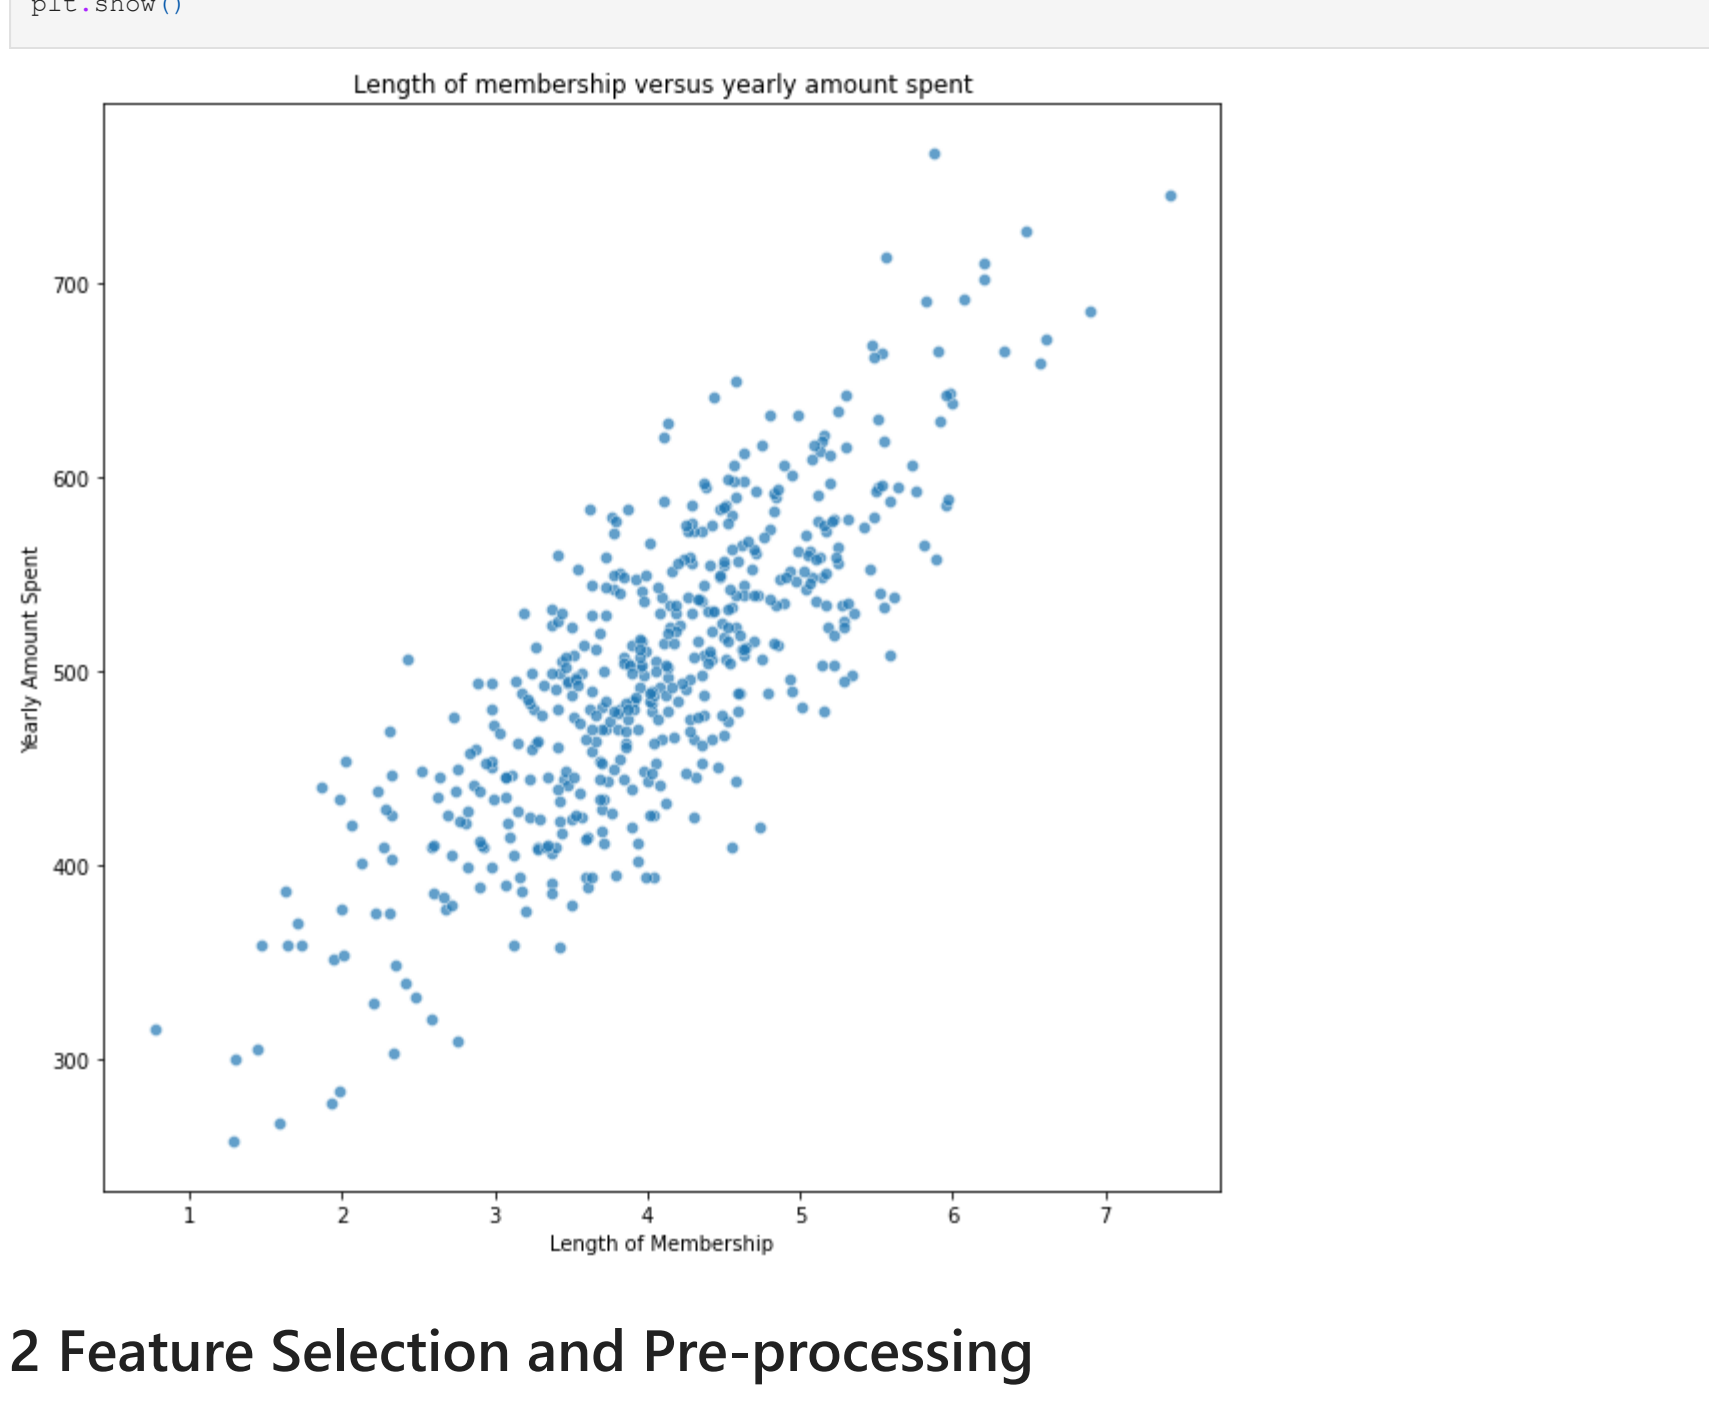


Length of Membership is most correlated with Yearly amount spent.



It is clear that Credit Card and Time on Website should be removed from the dataset as there is no correlation at all between the two. Average Session Length and Time on App have slight positive correlation but appear to be noise, but it is clear that Length of membership is the most related to yearly amount spent. Non Numerical Columns should also be removed as email cannot predict yearly amount spent. Physical address could if it was region based, but that is beyond the scope of this dataset as there are not enough samples to check if this persists so it also should be dropped.

Note: Time on Website would be dropped as it is in the instructions not to drop it.



2 Feature Selection and Pre-processing

Based on the EDA and null analysis, drop the unnecessary columns for the regression

```
In [13]: df_clean = df.drop(columns = ["Email", "Address", "Credit Card"])
```

3 X/Y and Training/Test Split

Use sklearn's train_test_split to split the data set into training and test sets. There should be 30% records in the test set. The random state should be 101. As we will be doing gradient descent as well as some other regression techniques, scaling the data set is important. So, use sklearn's StandardScaler for scaling the X of training and test sets. But don't do it for y(target) train and test. For help, you can see the answer for this question: <https://stackoverflow.com/questions/38780302/predicting-new-data-using-sklearn-after-standardizing-the-training-data>

```
In [14]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
X = df_clean.iloc[:, 0:4].values
y = df_clean.iloc[:, 4].values
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3, random_state=101)
```

4 Training Linear Model using SKLearn's LinearRegression

1. Train a linear model using Sklearn's LinearRegression (example in the linear regression slide/colab links in webcourses)
2. After training, show the coefficients and intercept
3. Predict for the test data
4. Generate a scatter plot that shows the Y test on x-axis and y predicted in y-axis
5. Use sklearn's metrics to print the value of MAE, MSE, RMSE, and R² (see documentation of sklearn's metrics)
6. Interpretation: Interpret the coefficient and which coefficient belongs to which feature and based on that explain any strategy that should help the business

```
In [15]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
reg=LinearRegression()
reg.fit(X_train, y_train)

Out[15]: LinearRegression()
```

```
In [16]: print("Best value of theta:", reg.intercept_, reg.coef_, sep="\n")

Best value of theta:
[500.5316775]
[25.76252659 38.32855202 0.19220992 61.17355707]
```

```
In [17]: y_pred = reg.predict(X_test)

Out[17]:
```



```
In [19]: from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import r2_score
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = mean_squared_error(y_test, y_pred, squared = False)
r2 = r2_score(y_test, y_pred)
print("MAE: {} \nMSE: {} \nRMSE: {} \nR_squared: {}".format(mae, mse, rmse, r2))

MAE: 7.228148653430837
MSE: 79.81305165979146
RMSE: 8.933815066978642
R_squared: 0.989046246741234
```

Coefficients are (25.76252659 38.32855202 0.19220992 61.17355707), which correspond to Avg. Session Length, Time on App, Time on Website, and Length of Membership in that order. As predicted, Length of membership has the biggest effect on the model while Time on Website has a insignificant effect. The other two have some effect, but much less than Length of membership.

5 Normal Equation

(while solving this, you might need to convert your dataframe into some different data structures such as to numpy(), might need to reshape, perform Transpose, add x0 columns, etc. I would recommend you to see the colab link I have shown you in the class and try to compare the shape of X and y and their data to get an idea during this process. Also, the code and discussion in the slide will help.)

1. Implement Normal Equation and find best_theta values based on the training set
2. Display the theta values. Are they very close to the sklearn's linear regression?
3. Prepare the test set before prediction
4. Perform Prediction for the test set
5. Generate a scatter plot that shows the Y test on x-axis and y predicted in y-axis
6. Use sklearn's metrics to print the value of MAE, MSE, RMSE, and R² (see documentation of sklearn's metrics)
7. Short Question: What are the benefits and the limitations of using batch gradient descent?

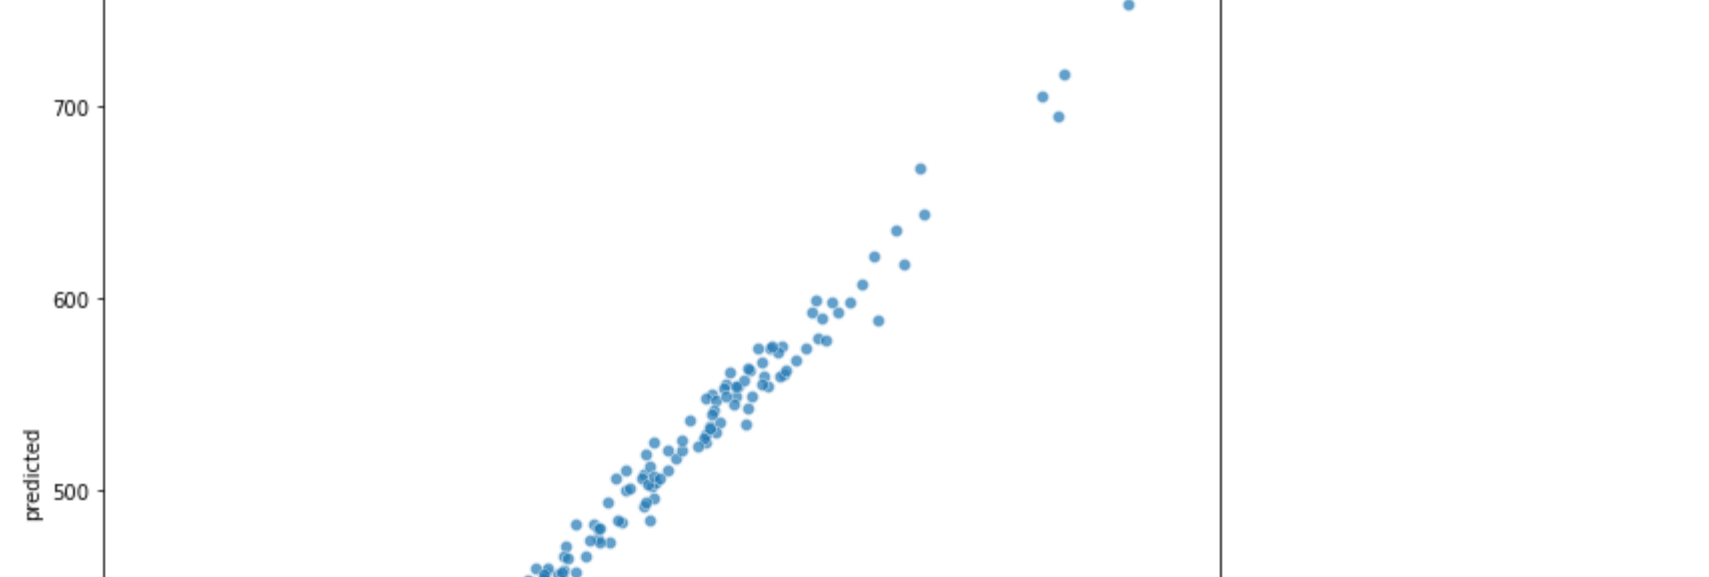
```
In [20]: X1_train = X_train[:,0]
X2_train = X_train[:,1]
X3_train = X_train[:,2]
X4_train = X_train[:,3]
X1_test = X_test[:,0]
X2_test = X_test[:,1]
X3_test = X_test[:,2]
X4_test = X_test[:,3]
m = len(X1_train)
n = len(X1_test)
X_train_bias = np.ones((m,1))
X_test_bias = np.ones((n,1))
X1_train = np.reshape(X1_train, (m,1))
X2_train = np.reshape(X2_train, (m,1))
X3_train = np.reshape(X3_train, (m,1))
X4_train = np.reshape(X4_train, (m,1))
X1_test = np.reshape(X1_test, (n,1))
X2_test = np.reshape(X2_test, (n,1))
X3_test = np.reshape(X3_test, (n,1))
X4_test = np.reshape(X4_test, (n,1))
XB_train = np.concatenate((X_train_bias, X1_train, X2_train, X3_train, X4_train), axis = 1)
XB_train_T = np.transpose(XB_train)
XB_test_T = np.transpose(XB_test)
XB_train_T.dot(XB_train) = XB_train_T.dot(XB_train)
XB_test_T.dot(XB_test) = XB_test_T.dot(XB_test)
train_theta = np.linalg.inv(XB_train_T.dot(XB_train).dot(XB_train_T.dot(y_train[:,0])))
print("Theta 1: {} \nTheta 2: {} \nTheta 3: {} \nTheta 4: {}".format(train_theta[0], train_theta[1], train_theta[2], train_theta[3]))

Theta 1: 500.53167757303
Theta 2: 25.76252659451966
Theta 3: 38.32855202493705
Theta 4: 0.1922099219945146
```

These values are very close except for theta 1, which is much higher than in the sklearn

```
In [21]: test_theta = np.linalg.inv(XB_test_T.dot(XB_test).dot(XB_test_T.dot(y_test[:,0])))
print("Theta 1: {} \nTheta 2: {} \nTheta 3: {} \nTheta 4: {}".format(test_theta[0], test_theta[1], test_theta[2], test_theta[3]))

Theta 1: 499.8932024429829
Theta 2: 24.90292432338286
Theta 3: 38.75765230036739
Theta 4: 0.774406826727872
```



```
In [23]: mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = mean_squared_error(y_test, y_pred, squared = False)
r2 = r2_score(y_test, y_pred)
print("MAE: {} \nMSE: {} \nRMSE: {} \nR_squared: {}".format(mae, mse, rmse, r2))

MAE: 7.09159957139494
MSE: 77.0614162693221
RMSE: 8.77846320658304
R_squared: 0.9893837013182971
```

A transpose dotted by A is O(N Cubed) and is expensive to run with large datasets.

6 Batch Gradient Descent

1. Implement Batch Gradient Descent based on the way we have learned in the class (See sample code form pdf). You can play with eta and n_iterations and should set to reasonable eta and number of iterations so that you can get the thetas close to Normal equation's theta
2. Display the theta values. Are they very close to the sklearn's linear regression?
3. Also, plot step number (in x-axis) against the cost/y-axis. See an example from this colab link: https://colab.research.google.com/drive/19_UoIFBx-n0Dfs7Pw7Z8BgVw5yQLe?usp=sharing (Links to an external site.)
4. Perform Prediction for the test set
5. Generate a scatter plot that shows the Y test on the x-axis and y predicted in the y-axis
6. Use sklearn's metrics to print the value of MAE, MSE, RMSE, and R² (see documentation of sklearn's metrics)
7. Short Question: How do derivatives help in the process of gradient descent?

```
In [24]: cost_list = []
epoch_list = []
predicted_list = []

eta = 0.005 # learning rate
n_iterations = 500

theta = np.random.randn(5,1) # random initialization

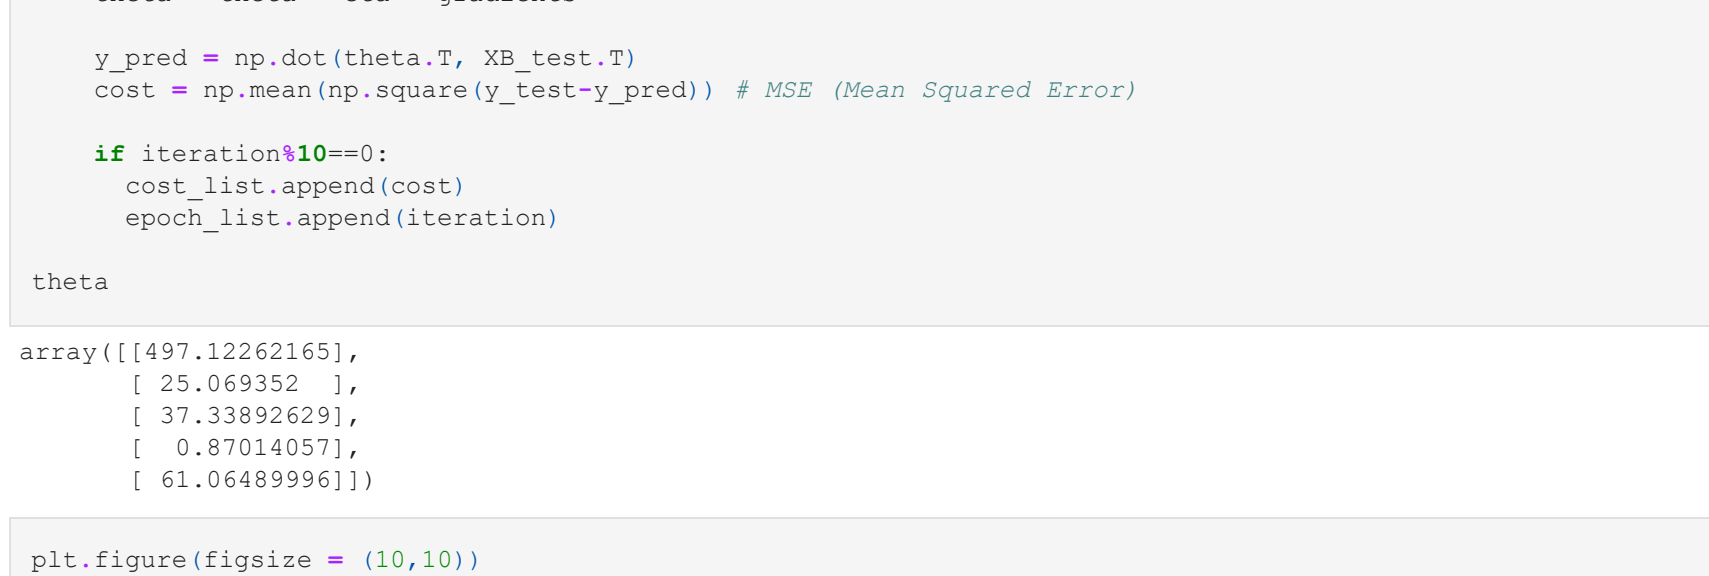
for iteration in range(n_iterations):
    gradients = 2/m * XB_train.T.dot(XB_train.dot(theta) - y_train)
    theta = theta - eta * gradients

    y_pred = np.dot(theta, XB_test.T)
    cost = np.mean(np.square(y_test-y_pred)) # MSE (Mean Squared Error)

    if iteration%10==0:
        cost_list.append(cost)
        epoch_list.append(iteration)

    theta

Out[24]: array([[497.12262165],
       [25.069352 ],
       [37.7921462],
       [ 0.87014055],
       [ 61.06489996]])
```



```
In [27]: mae = mean_absolute_error(y_test, y_pred[0,:])
mse = mean_squared_error(y_test, y_pred[0,:])
rmse = mean_squared_error(y_test, y_pred[0,:], squared = False)
r2 = r2_score(y_test, y_pred[0,:])
print("MAE: {} \nMSE: {} \nRMSE: {} \nR_squared: {}".format(mae, mse, rmse, r2))

MAE: 7.338263007748179
MSE: 89.560632023617
RMSE: 9.463860580209845
R_squared: 0.987795219070836
```

Derivatives help minimize the cost function. More stable convergence and error gradient than Stochastic Gradient descent. Can converge at local minima and saddle points

7 Stochastic Gradient Descent

1. Implement Stochastic Gradient Descent and train our data set. You must have to use learning_schedule (see example code in pdf as well as the colab link I have shared in #6 above). The parameters should be reasonable and the theta values should be very close to the normal equation
2. Display the theta values. Are they very close to the sklearn's linear regression?
3. Also, plot step number (in x-axis) against cost/y-axis. See an example from this colab link: https://colab.research.google.com/drive/19_UoIFBx-n0Dfs7Pw7Z8BgVw5yQLe?usp=sharing (Links to an external site.)
4. Perform Prediction for the test set
5. Generate a scatter plot that shows the Y test on the x-axis and y predicted in the y-axis
6. Use sklearn's metrics to print the value of MAE, MSE, RMSE, and R² (see documentation of sklearn's metrics)
7. Short Question: What are the benefits and the limitations of using batch gradient descent?


```
[28]: cost_list = []
      epoch_list = []
      predicted_list = []

      n_epochs = 200
      t0, t1 = 500, 5000 # learning schedule hyperparameters

      def learning_schedule(t):
          return t0 / (t + t1)

      theta = np.random.randn(5,1) # random initialization
      l=0
      for epoch in range(n_epochs):
          if epoch==0:
              i = 0
              random_index = np.random.randint(m)
              #print(random_index)
              xi = XB_train[random_index:random_index+1]
              yi = y_train[random_index:random_index+1]
              gradients = 2 * xi.T.dot(xi.dot(theta) - yi)
              eta = learning_schedule(epoch * m + i)
              theta = theta - eta * gradients
              i += 1

              y_pred = np.dot(theta.T, XB_test.T)
              cost = np.mean(np.square(y_test-y_pred)) # MSE (Mean Squared Error)

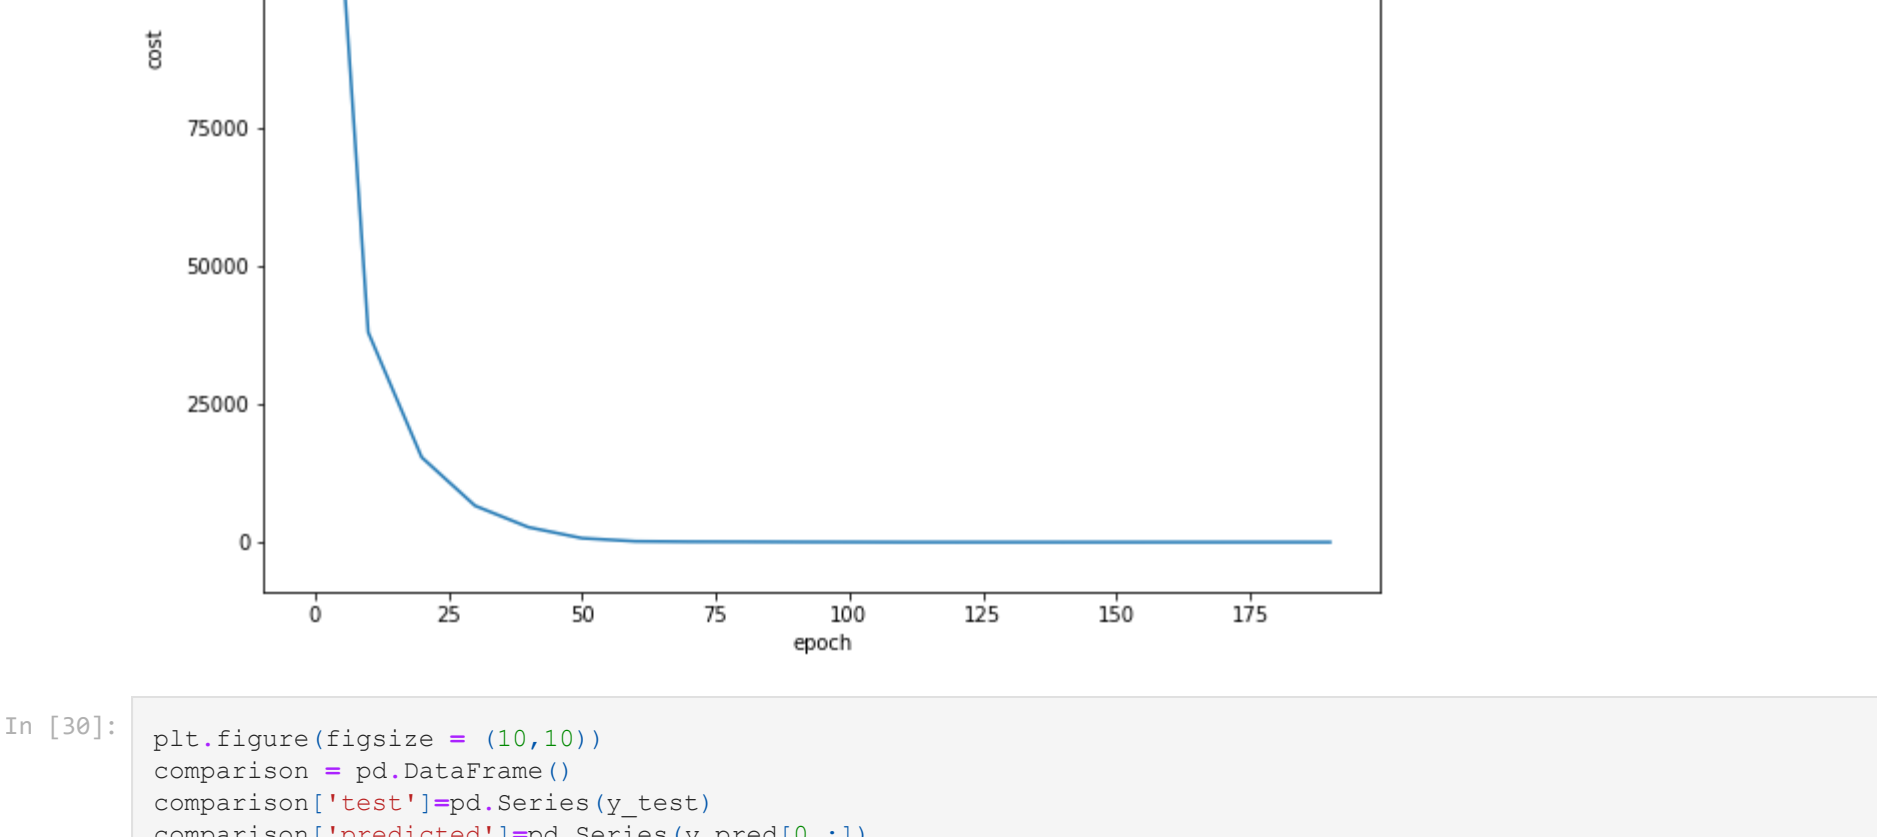
              if epoch==0:
                  cost_list.append(cost)
                  epoch_list.append(epoch)

      theta
```

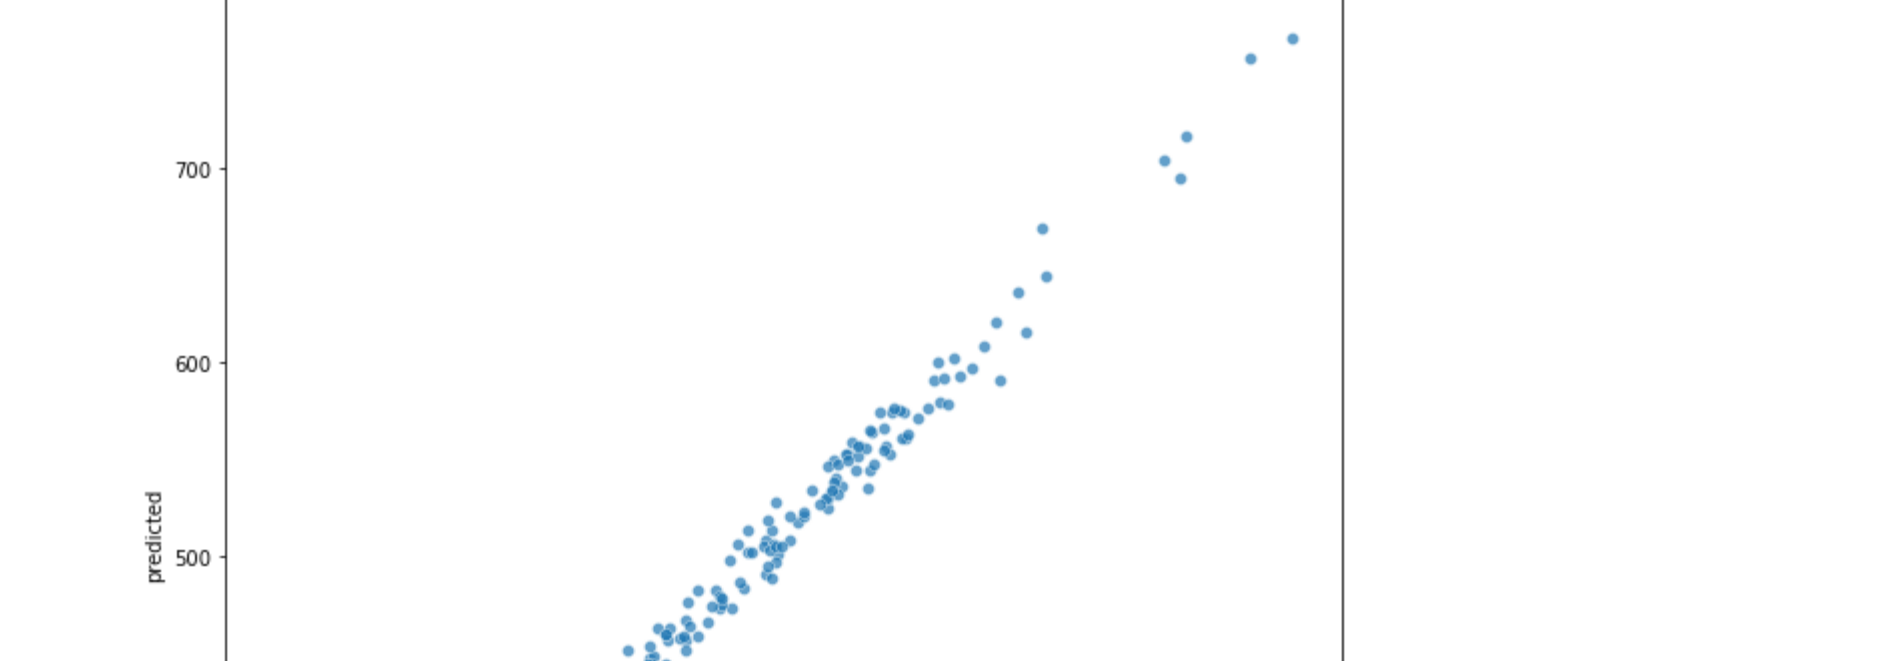
```
Out[28]: array([[ 1.00120749e+02],
        [ 2.60568941e+01],
        [ 3.84859364e+01],
        [ 1.92987914e+01],
        [ 4.13292936e+01]])
```

```
In [29]: plt.figure(figsize = (10,10))
          plt.xlabel("epoch")
          plt.ylabel("cost")
          plt.plot(epoch_list,cost_list)
```

```
Out[29]: (<matplotlib.lines.Line2D at 0x24011899700>)
```



```
In [30]: plt.figure(figsize = (10,10))
          comparison = pd.DataFrame()
          comparison["test"] = pd.Series(y_test)
          comparison["predicted"] = pd.Series(y_pred(0,:))
          sns.scatterplot(data = comparison, x = "test", y = "predicted", alpha = 0.7).set(title = "Test versus Predicted Values")
          plt.show()
```



```
In [31]: mae = mean_absolute_error(y_test, y_pred(0,:))
          mse = mean_squared_error(y_test, y_pred(0,:))
          rmse = mean_squared_error(y_test, y_pred(0,:), squared = False)
          r2 = r2_score(y_test, y_pred(0,:))
          print("MAE: {} \nMSE: {} \nRMSE: {} \nR_squared: {}".format(mae, mse, rmse, r2))

MAE: 7.28177320718996
MSE: 81.2893398060397
RMSE: 9.016093378322992
R_squared: 0.98801624677553
```

For larger datasets, it can converge faster as it causes updates to the parameters more frequently. Due to frequent updates, the steps taken towards the minima are very noisy. This can often lead the gradient descent into other directions.

8 SGDRegressor from sklearn

1. Use sklearn's SGDRegressor to train a model for our data set. Put a reasonable iteration and tolerance and learning steps so that we can get coefficients close to normal equation
2. Display the theta values. Are they very close to sklearn's linear regression?
3. Predict for the test data
4. Generate a scatter plot that shows the Y test on the x-axis and y predicted in the y-axis
5. Use sklearn's metrics to print the value of MAE, MSE, RMSE, and R^2 (see documentation of sklearn's metrics)

```
In [32]: from sklearn.linear_model import SGDRegressor
          sgd_reg = SGDRegressor(max_iter=1000, tol=1e-3, penalty=None, eta=0.1)
          sgd_reg.fit(X_train, y_train.ravel())
          sgd_reg.intercept_, sgd_reg.coef_
```

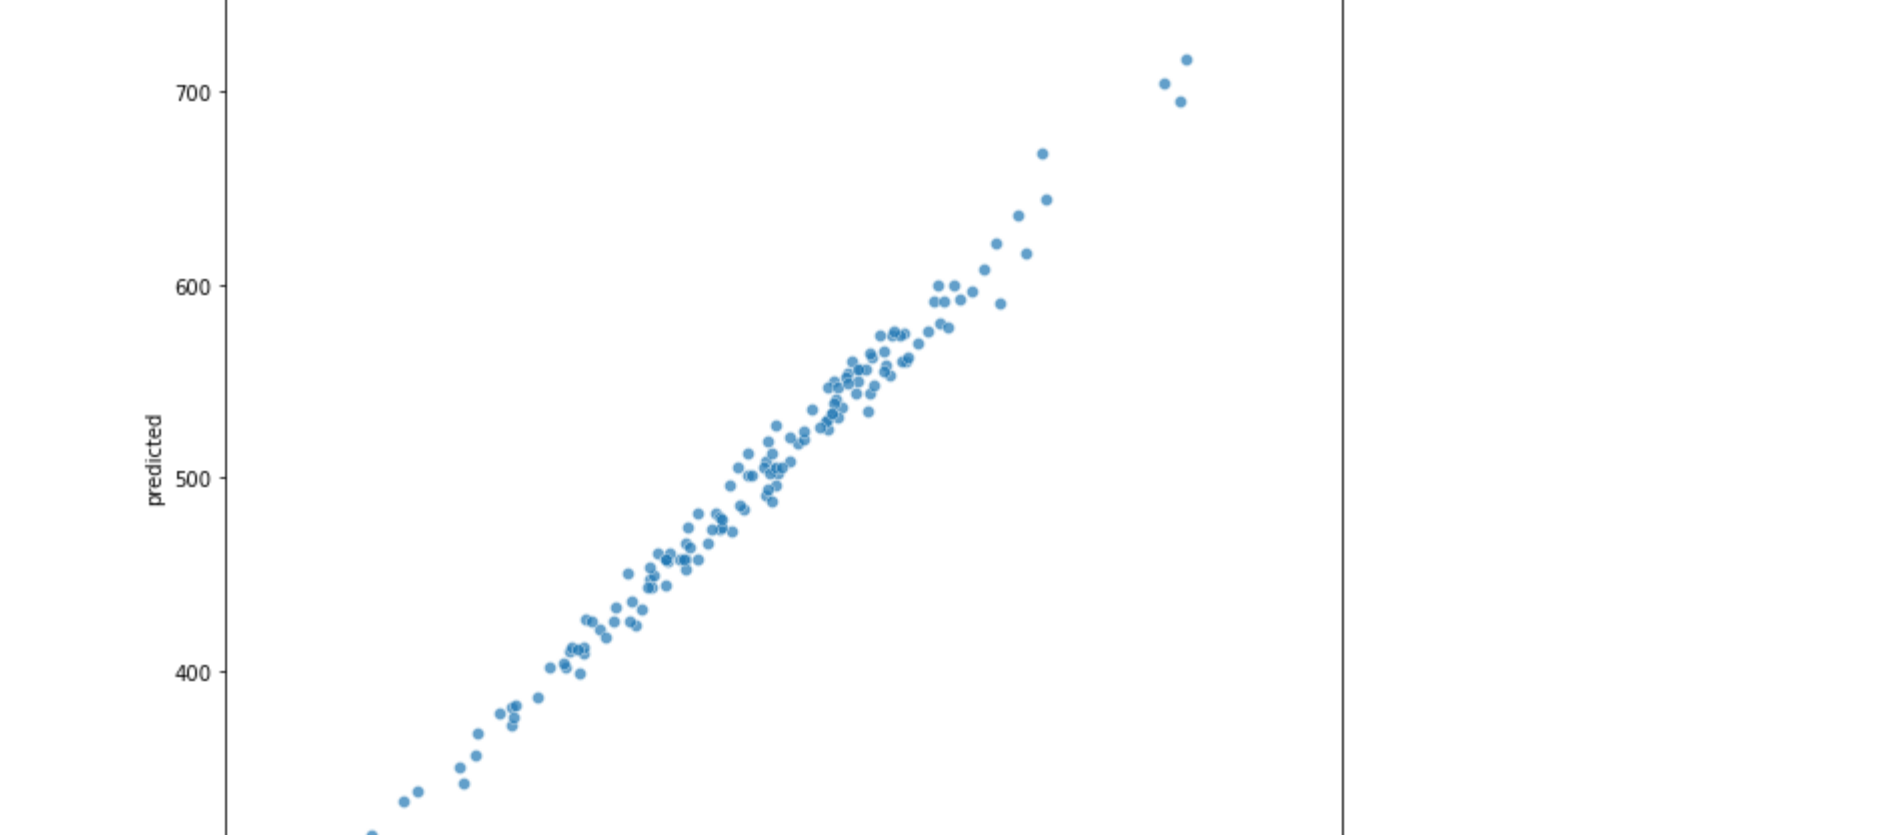
```
Out[32]: (array([300.35067937],
        [25.5930067366, 36.67331153, -0.09914352, 61.48881844]))

They are close to Sklearn's linear regression.
```

```
In [33]: y_pred = sgd_reg.predict(X_test)
```

```
In [34]: plt.figure(figsize = (10,10))
          comparison = pd.DataFrame()
          comparison["test"] = pd.Series(y_test)
          comparison["predicted"] = pd.Series(y_pred)
```

```
sns.scatterplot(data = comparison, x = "test", y = "predicted", alpha = 0.7).set(title = "Test versus Predicted Values")
plt.show()
```



```
In [35]: mae = mean_absolute_error(y_test, y_pred)
          mse = mean_squared_error(y_test, y_pred)
          rmse = mean_squared_error(y_test, y_pred, squared = False)
          r2 = r2_score(y_test, y_pred)
          print("MAE: {} \nMSE: {} \nRMSE: {} \nR_squared: {}".format(mae, mse, rmse, r2))

MAE: 7.209153003234533
MSE: 79.19536986979392
RMSE: 8.899178050177383
R_squared: 0.9890897190513886
```

9 Mini-batch Gradient Descent

Briefly explain how mini-batch can overcome the limitations of Batch gradient descent and SGD.

Easily fits in the memory. It is computationally efficient. Benefit from vectorization. If stuck in local minimums, some noisy steps can lead the way out of them. Average of the training samples produces stable error gradients and convergence.

10 Polynomial of degree 2

1. Use sklearn's Polynomial features to degree = 2 on our training and test set
2. Use linearRegression on the new polynomial features
3. Predict for test set
4. Generate a scatter plot that shows the Y test on the x-axis and y predicted in the y-axis
5. Use sklearn's metrics to print the value of MAE, MSE, RMSE, and R^2 (see documentation of sklearn's metrics)

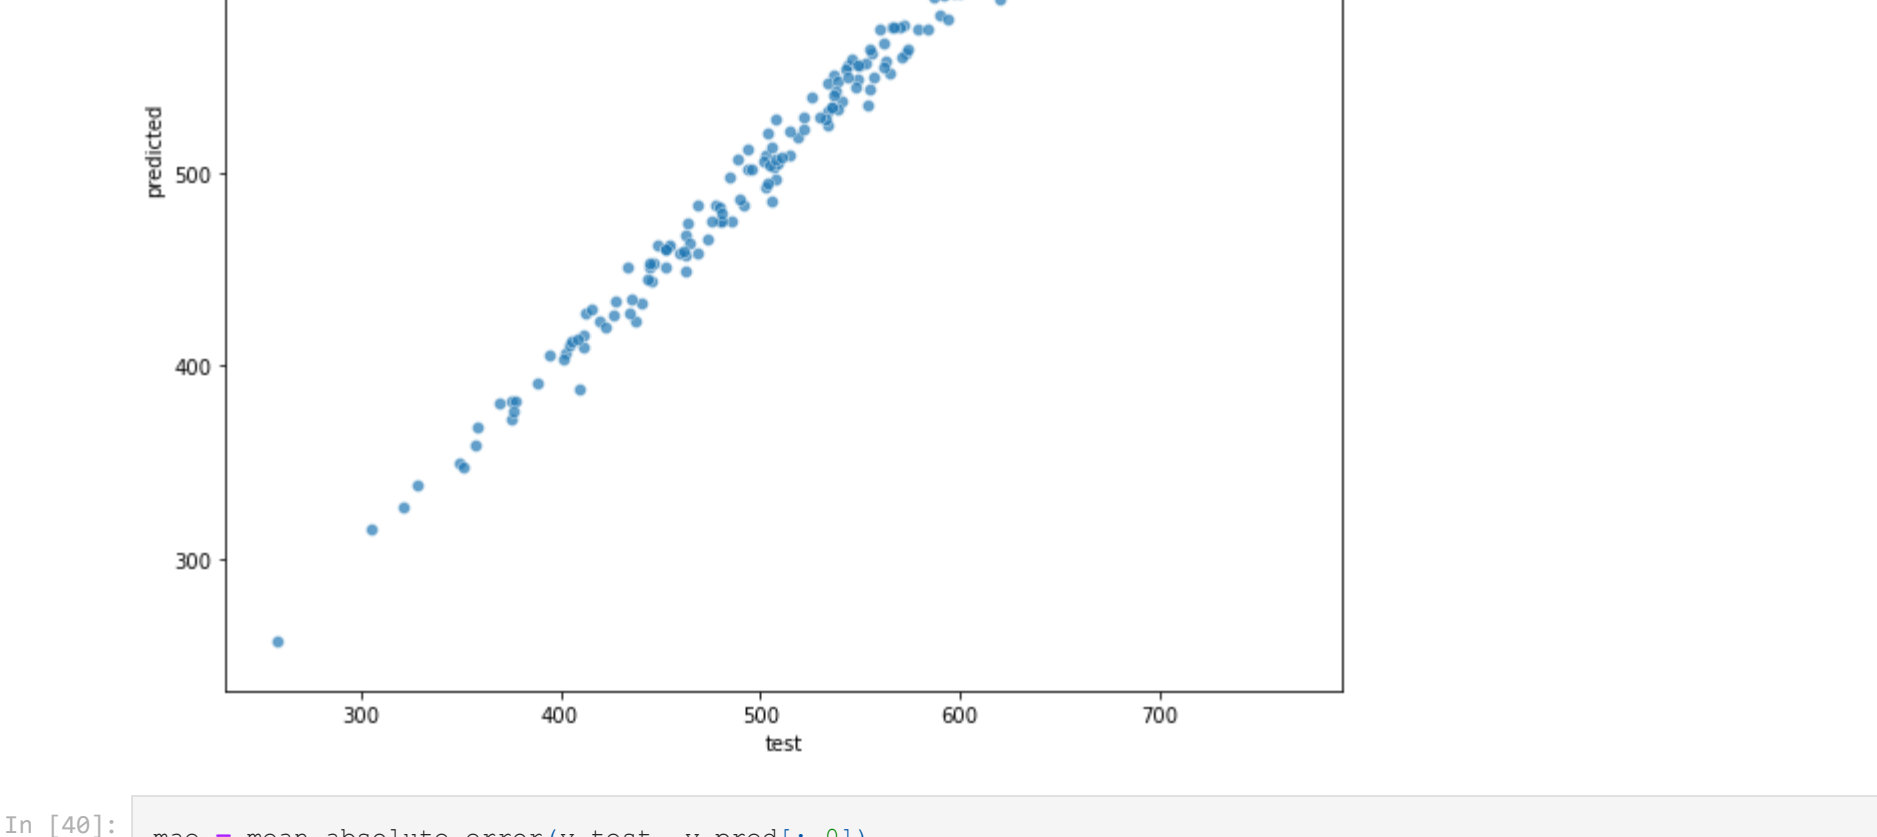
```
In [36]: from sklearn.preprocessing import PolynomialFeatures
          poly = PolynomialFeatures(2)
          poly_X_train = poly.fit_transform(X_train)
          poly_X_test = poly.fit_transform(X_test)
```

```
In [37]: reg=LinearRegression()
          reg.fit(poly_X_train, y_train)
```

```
Out[37]: LinearRegression()
```

```
In [38]: y_pred = reg.predict(poly_X_test)
```

```
In [39]: plt.figure(figsize = (10,10))
          comparison = pd.DataFrame()
          comparison["test"] = pd.Series(y_test)
          comparison["predicted"] = pd.Series(y_pred(0,:))
          comparison["predicted"] = pd.Series(y_pred(0,:))
          sns.scatterplot(data = comparison, x = "test", y = "predicted", alpha = 0.7).set(title = "Test versus Predicted Values")
          plt.show()
```



```
In [40]: mae = mean_absolute_error(y_test, y_pred(0,:))
          mse = mean_squared_error(y_test, y_pred(0,:))
          rmse = mean_squared_error(y_test, y_pred(0,:), squared = False)
          r2 = r2_score(y_test, y_pred(0,:))
          print("MAE: {} \nMSE: {} \nRMSE: {} \nR_squared: {}".format(mae, mse, rmse, r2))

MAE: 7.435295703827286
MSE: 85.34745505512038
RMSE: 9.238368636024621
R_squared: 0.9882481824247091
```

11 Polynomial of degree 3

1. Use sklearn's Polynomial features to degree = 3 on our training and test set
2. Use linearRegression on the new polynomial features
3. Predict for test set
4. Generate a scatter plot that shows the Y test on the x-axis and y predicted in the y-axis
5. Use sklearn's metrics to print the value of MAE, MSE, RMSE, and R^2 (see documentation of sklearn's metrics)

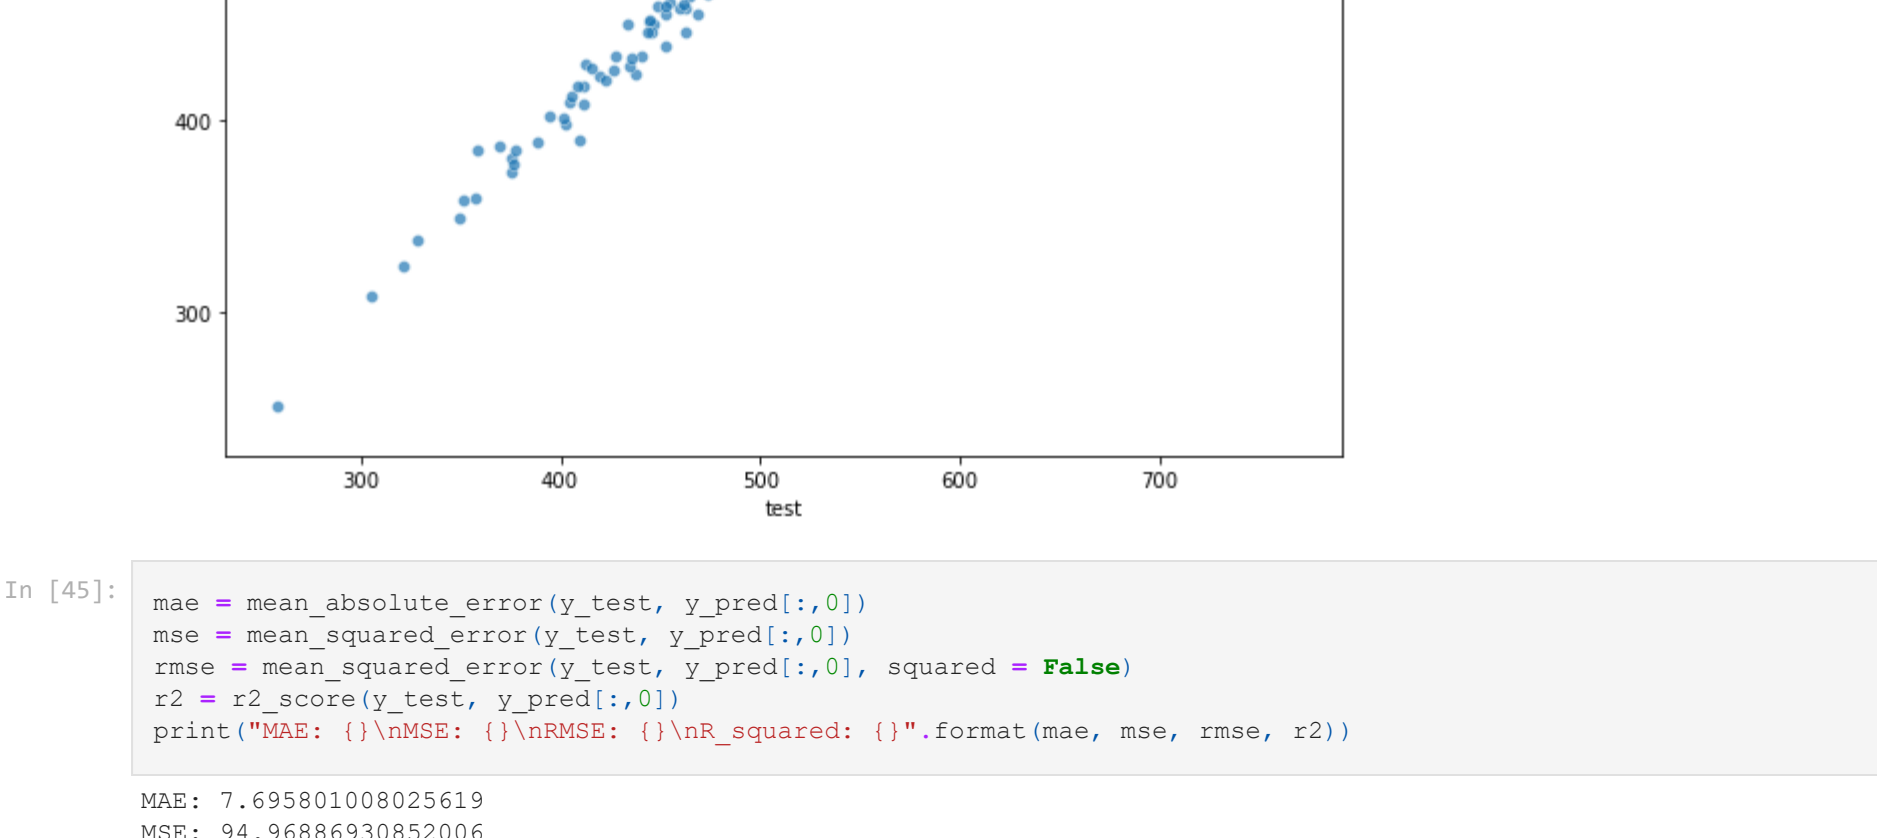
```
In [41]: poly = PolynomialFeatures(3)
          poly_X_train = poly.fit_transform(X_train)
          poly_X_test = poly.fit_transform(X_test)
```

```
In [42]: reg=LinearRegression()
          reg.fit(poly_X_train, y_train)
```

```
Out[42]: LinearRegression()
```

```
In [43]: y_pred = reg.predict(poly_X_test)
```

```
In [44]: plt.figure(figsize = (10,10))
          comparison = pd.DataFrame()
          comparison["test"] = pd.Series(y_test)
          comparison["predicted"] = pd.Series(y_pred(0,:))
          comparison["predicted"] = pd.Series(y_pred(0,:))
          sns.scatterplot(data = comparison, x = "test", y = "predicted", alpha = 0.7).set(title = "Test versus Predicted Values")
          plt.show()
```



```
In [45]: mae = mean_absolute_error(y_test, y_pred(0,:))
          mse = mean_squared_error(y_test, y_pred(0,:))
          rmse = mean_squared_error(y_test, y_pred(0,:), squared = False)
          r2 = r2_score(y_test, y_pred(0,:))
          print("MAE: {} \nMSE: {} \nRMSE: {} \nR_squared: {}".format(mae, mse, rmse, r2))

MAE: 7.65950108025619
MSE: 94.96886930852006
RMSE: 9.745197243181097
R_squared: 0.9863656966965793
```

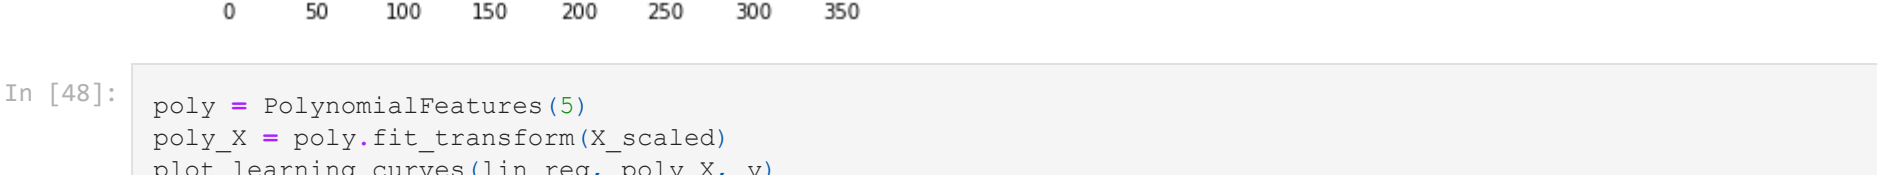
12 Learning Curve

1. Generate learning curve with linearRegression
2. Generate learning curve with polynomial regression with degree = 5
3. Interpret the result

```
In [46]: def plot_learning_curves(model, X, y):
          X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=101)
          for m in range(1, len(X_train)):
              model.fit(X_train[:m], y_train[:m])
              y_train_predict = model.predict(X_train[:m])
              y_test_pred = model.predict(X_test)
              train_errors.append(mean_squared_error(y_train[:m], y_train_predict))
              pred_errors.append(mean_squared_error(y_test, y_test_pred))
              plt.plot(np.sqrt(train_errors), "r-+", linewidth=2, label="train")
              plt.plot(np.sqrt(pred_errors), "b-", linewidth=3, label="test")
              print("Train Error: {}".format(train_errors[348]), pred_errors[348])
```

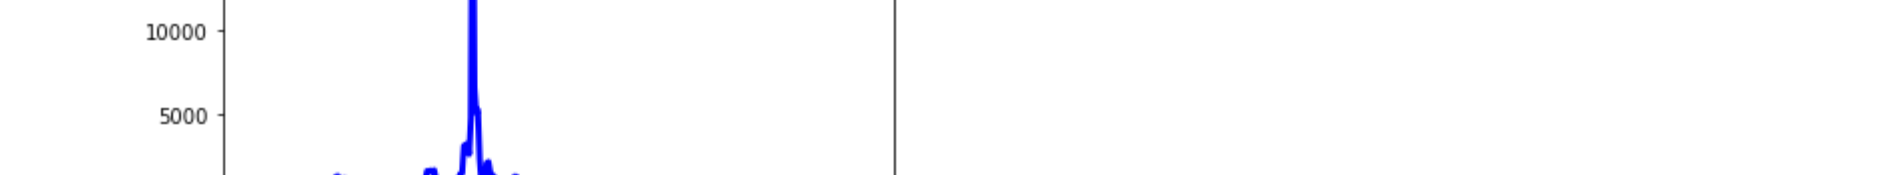
```
In [47]: lin_reg = LinearRegression()
          plot_learning_curves(lin_reg, X_scaled, y)
```

```
Train Error: 107.44878264719886
Val Error: 79.85787475413264
```



```
In [48]: poly = PolynomialFeatures(5)
          poly_X = poly.fit_transform(X_scaled)
          plot_learning_curves(lin_reg, poly_X, y)
```

```
Train Error: 73.15139658106436
Val Error: 6395.80957778272
```



The model with data fitted to a 5 feature Polynomial is showing considerable signs of overfitting, with a much higher validation error than training error.

13 Regularization

Explain the purpose of regularization. For the following Regularization method, use the polynomial degree 3 data set

Since the polynomial dataset is prone to overfitting, regularization can help this issue by constraining the weights of the model.

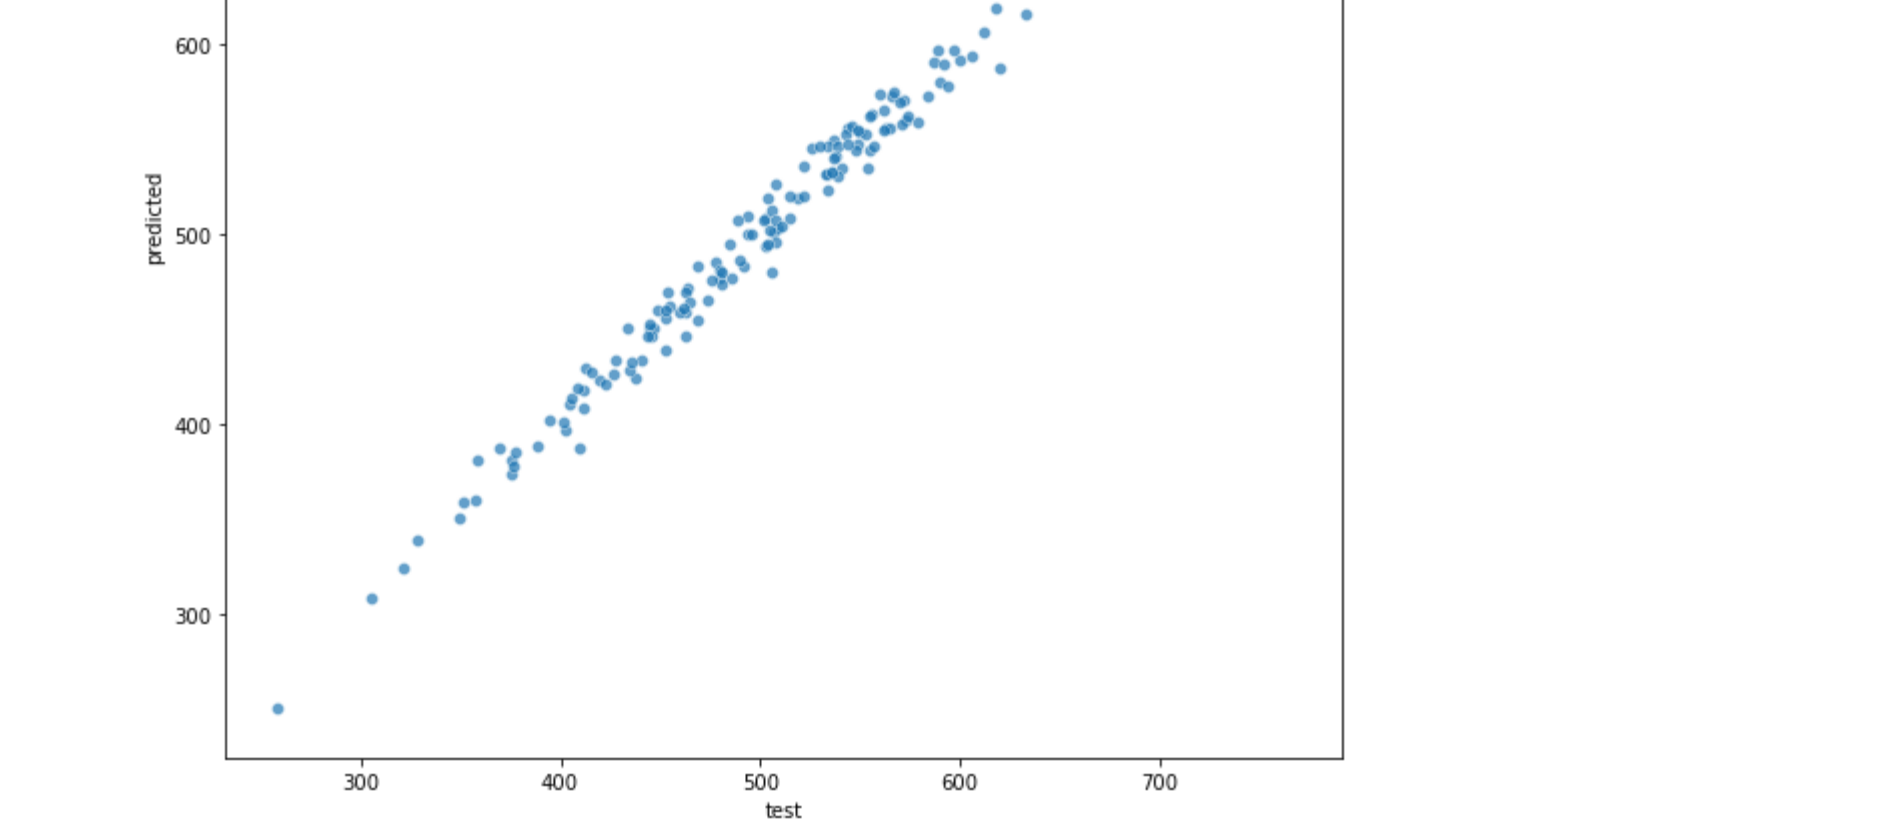
14 Ridge Regression

1. Use sklearn's Ridge to train the data set (use the polynomial degree 3 data set)
2. Predict for test set
3. Generate a scatter plot that shows the Y test on the x-axis and y predicted in the y-axis
4. Use sklearn's metrics to print the value of MAE, MSE, RMSE, and R^2 (see documentation of sklearn's metrics)

```
In [49]: poly = PolynomialFeatures(3)
          poly_X_train = poly.fit_transform(X_train)
          poly_X_test = poly.fit_transform(X_test)
```

```
In [50]: from sklearn.linear_model import Ridge
          ridge_reg = Ridge(alpha=1, solver="cholesky")
          ridge_reg.fit(poly_X_train, y_train)
          y_pred = ridge_reg.predict(poly_X_test)
```

```
In [51]: plt.figure(figsize = (10,10))
          comparison = pd.DataFrame()
          comparison["test"] = pd.Series(y_test)
          comparison["predicted"] = pd.Series(y_pred)
          comparison["predicted"] = pd.Series(y_pred)
          sns.scatterplot(data = comparison, x = "test", y = "predicted", alpha = 0.7).set(title = "Test versus Predicted Values")
          plt.show()
```



```
In [52]: mae = mean_absolute_error(y_test, y_pred(0,:))
          mse = mean_squared_error(y_test, y_pred(0,:))
          rmse = mean_squared_error(y_test, y_pred(0,:), squared = False)
          r2 = r2_score(y_test, y_pred(0,:))
          print("MAE: {} \nMSE: {} \nRMSE: {} \nR_squared: {}".format(mae, mse, rmse, r2))
```

```
MAE: 7.855355046312291
MSE: 97.6595914733688
RMSE: 9.88228673245364
R_squared: 0.986460117084953
```

15 SGDRegressor for Ridge

1. Use sklearn's SGDRegressor for Ridge Regression
2. Predict for test set
3. Generate a scatter plot that shows the Y test on the x-axis and y predicted in the y-axis
4. Use sklearn's metrics to print the value of MAE, MSE, RMSE, and R^2 (see documentation of sklearn's metrics)

```
In [53]: from sklearn.linear_model import SGDRegressor
          sgd_reg = SGDRegressor(penalty="l2")
          sgd_reg.fit(poly_X_train, y_train.ravel())
          y_pred = sgd_reg.predict(poly_X_test)
```

```
In [54]: plt.figure(figsize = (10,10))
          comparison = pd.DataFrame()
          comparison["test"] = pd.Series(y_test)
          comparison["predicted"] = pd.Series(y_pred)
          comparison["predicted"] = pd.Series(y_pred)
          sns.scatterplot(data = comparison, x = "test", y = "predicted", alpha = 0.7).set(title = "Test versus Predicted Values")
          plt.show()
```



```
In [55]: mae = mean_absolute_error(y_test, y_pred)
          mse = mean_squared_error(y_test, y_pred)
          rmse = mean_squared_error(y_test, y_pred, squared = False)
          r2 = r2_score(y_test, y_pred)
          print("MAE: {} \nMSE: {} \nRMSE: {} \nR_squared: {}".format(mae, mse, rmse, r2))
```

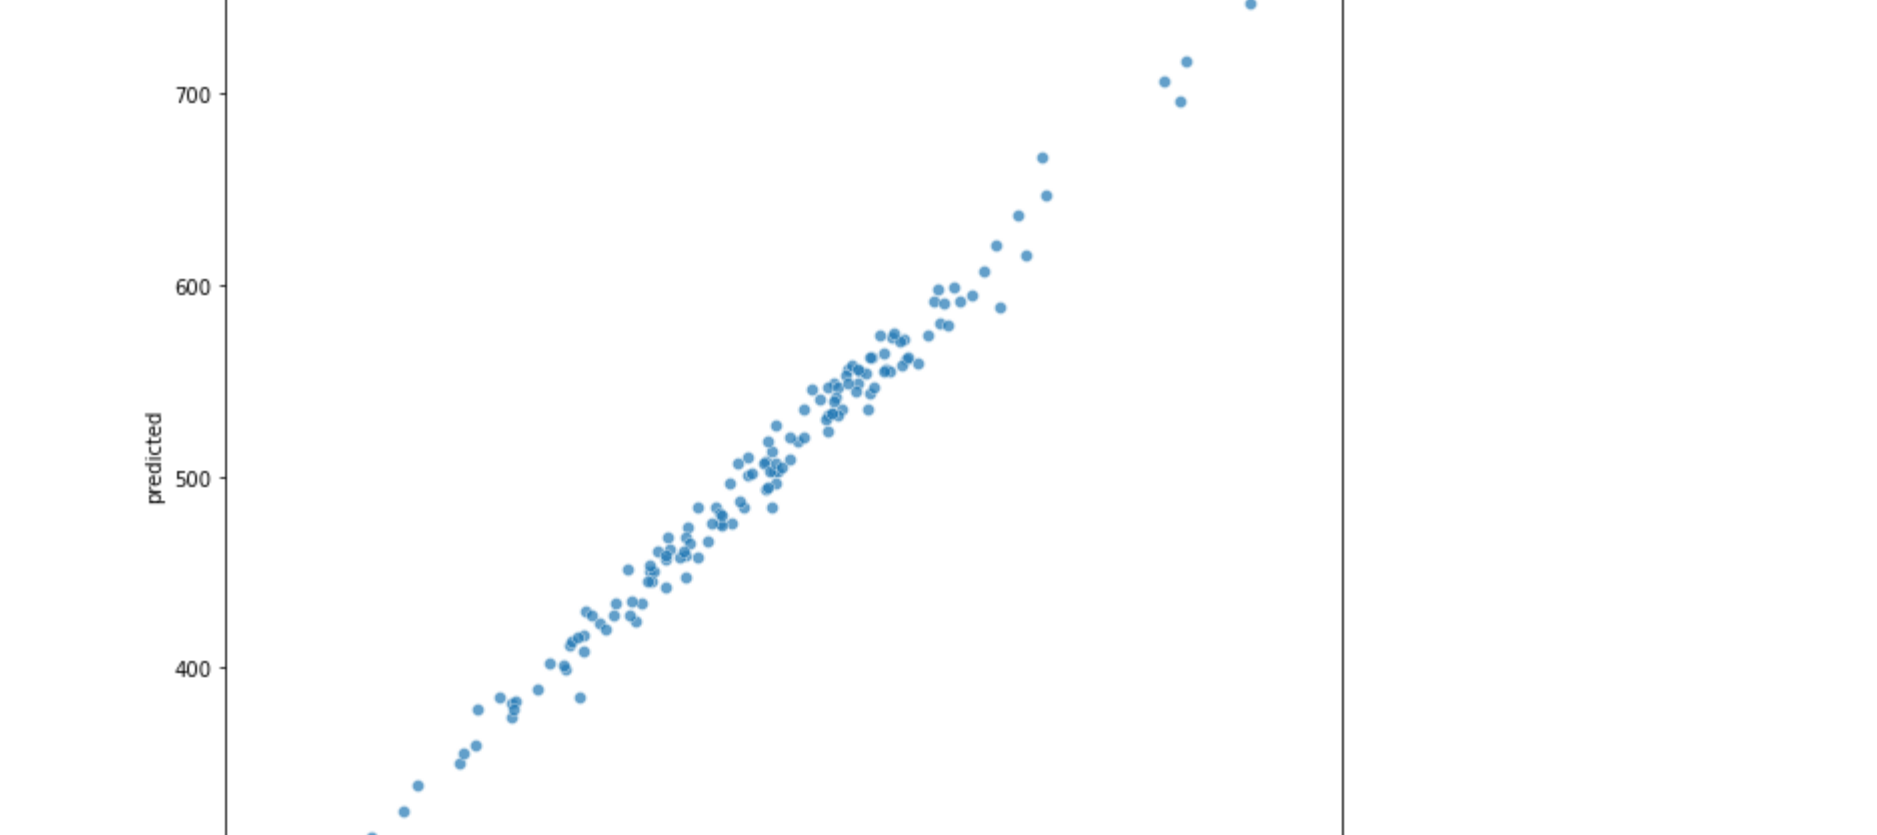
```
MAE: 7.679627114794209
MSE: 2126788.367467365
RMSE: 1458.351243868279
R_squared: -25129591421167556
```

16 Lasso Regression

1. Use sklearn's Lasso
2. Predict for test set
3. Generate a scatter plot that shows the Y test on the x-axis and y predicted in the y-axis
4. Use sklearn's metrics to print the value of MAE, MSE, RMSE, and R^2 (see documentation of sklearn's metrics)
5. How Lasso performs the regularization and how does that affect the thetas?

```
In [56]: from sklearn.linear_model import Lasso
          lasso_reg = Lasso(alpha = 0.1)
          lasso_reg.fit(poly_X_train, y_train)
          y_pred = lasso_reg.predict(poly_X_test)
```

```
In [57]: plt.figure(figsize = (10,10))
          comparison = pd.DataFrame()
          comparison["test"] = pd.Series(y_test)
          comparison["predicted"] = pd.Series(y_pred)
          comparison["predicted"] = pd.Series(y_pred)
          sns.scatterplot(data = comparison, x = "test", y = "predicted", alpha = 0.7).set(title = "Test versus Predicted Values")
          plt.show()
```



```
In [58]: mae = mean_absolute_error(y_test, y_pred)
          mse = mean_squared_error(y_test, y_pred)
          rmse = mean_squared_error(y_test, y_pred, squared = False)
          r2 = r2_score(y_test, y_pred)
          print("MAE: {} \nMSE: {} \nRMSE: {} \nR_squared: {}".format(mae, mse, rmse, r2))
```

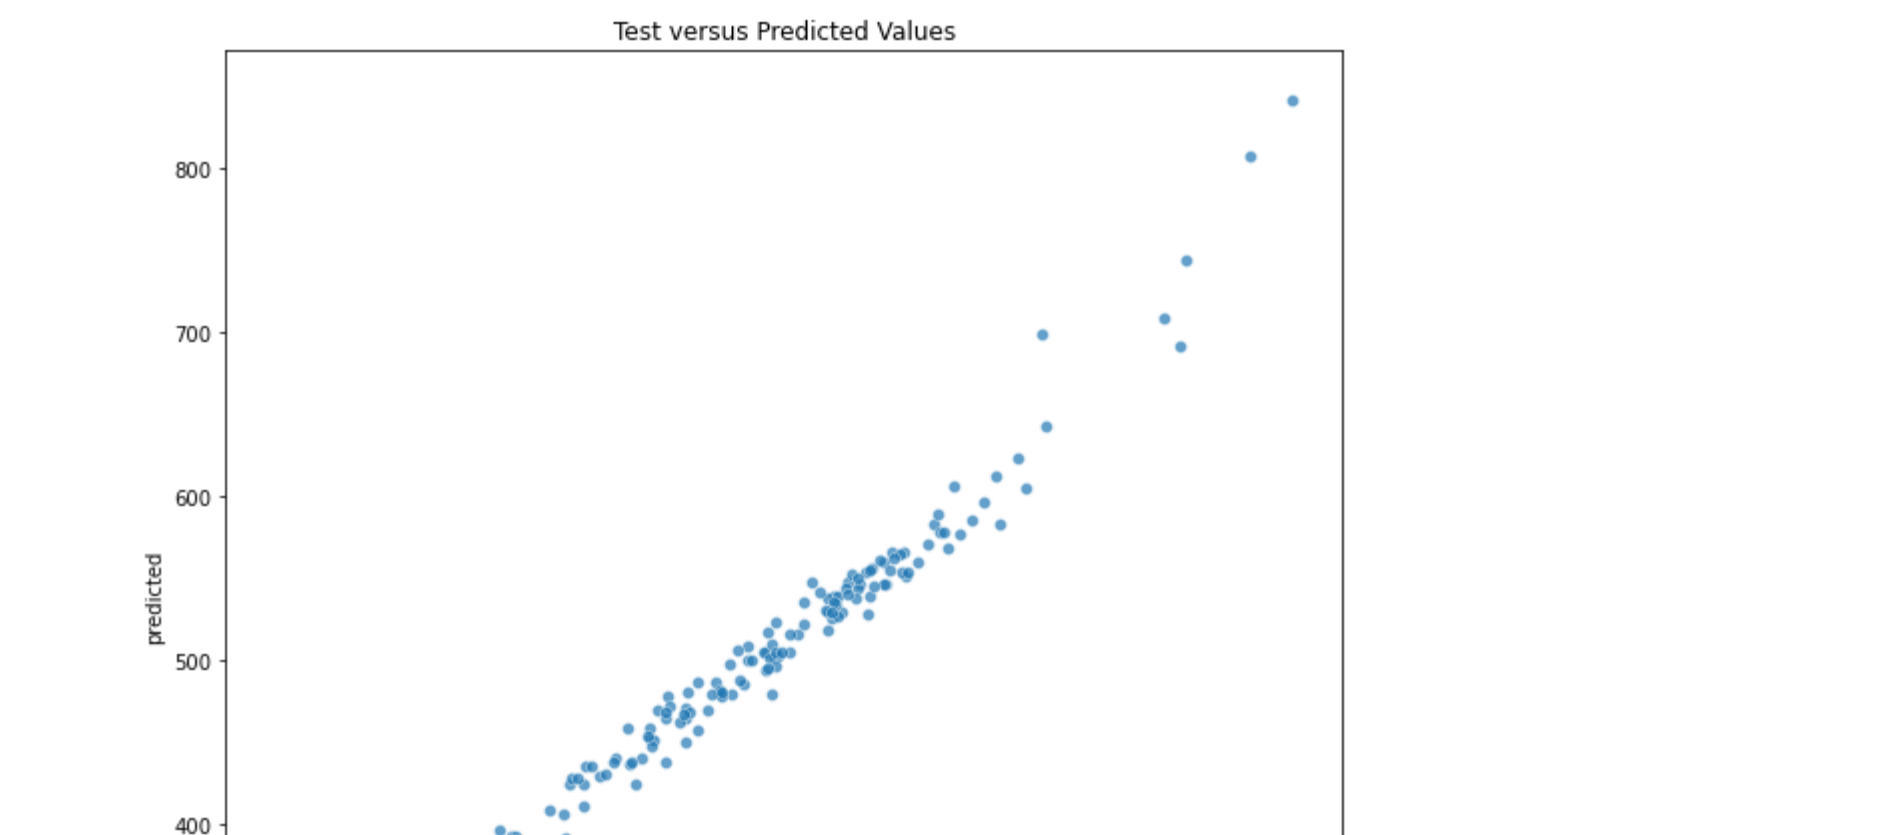
```
MAE: 7.679627114794209
MSE: 93.21215816943548
RMSE: 9.65464438375054
R_squared: 0.9891588471769
```

17 Elastic Net

1. Use sklearn's ElasticNet
2. Predict for test set
3. Generate a scatter plot that shows the Y test on x axis and y predicted in y axis
4. Use sklearn's metrics to print the value of MAE, MSE, RMSE, and R^2 (see documentation of sklearn's metrics)
5. How ElasticNet different compared to Lasso and RIDGE perform the regularization and how does that affect the thetas?

```
In [59]: from sklearn.linear_model import ElasticNet
          elastic_net = ElasticNet(alpha=0.1, l1_ratio=0.5)
          elastic_net.fit(poly_X_train, y_train)
          y_pred = elastic_net.predict(poly_X_test)
```

```
In [60]: plt.figure(figsize = (10,10))
          comparison = pd.DataFrame()
          comparison["test"] = pd.Series(y_test)
          comparison["predicted"] = pd.Series(y_pred)
          comparison["predicted"] = pd.Series(y_pred)
          sns.scatterplot(data = comparison, x = "test", y = "predicted", alpha = 0.7).set(title = "Test versus Predicted Values")
          plt.show()
```



```
In [61]: mae = mean_absolute_error(y_test, y_pred)
          mse = mean_squared_error(y_test, y_pred)
          rmse = mean_squared_error(y_test, y_pred, squared = False)
          r2 = r2_score(y_test, y_pred)
          print("MAE: {} \nMSE: {} \nRMSE: {} \nR_squared: {}".format(mae, mse, rmse, r2))
```

Elastic net is a mix of Ridge and Lasso regularization terms with a mix ratio of α . In general, Elastic Net is preferred over Lasso since Lasso may behave erratically when the number of features is greater than the number of training instances or when several features are strongly correlated.

18 Bonus Question:

In most of the above cases, for example, LinearRegression of sklearn, (Q4 above), we have used scaled data set for training. However, in a real-life scenario, you would like to predict the yearly amount spent for a new instance. The real data will not be scaled. How would you use the model for this case to predict this instance? (35.49726772511229,12.655651149166752,39.57766801952616,4.082620632952961) = ? \

Write necessary code so that it will predict a reasonable value for the amount spent. This is very close to our first training record.

```
In [62]: reg=LinearRegression()
          reg.fit(X_train, y_train)
          X_test = scaler.transform(np.expand_dims(np.array([35.49726772511229,12.655651149166752,39.57766801952616,4.082620632952961]),axis=0))
          y_pred = reg.predict(X_test)
```